

Writing Against the Machine: Computational Authorship and Historical Writing

CHRISTOPHER GERTEIS 

SOAS University of London

Abstract

Historians generate knowledge through the labour of composition – through the friction between interpretation and evidence that makes claims open to scrutiny and challenge. This essay argues that when composition is bypassed, that structure disappears. Generative AI raises this issue in urgent fashion. Current large language models produce what the essay terms ‘stochastic history’: prose that replicates the surface forms of historical explanation without enacting the disciplinary reasoning behind them. Such prose flattens temporal complexity into chronological adjacency, inherits narrative patterns without deliberating over them and reproduces hegemonic framings without the mechanisms – archival friction, peer contestation and historiographical consciousness – through which the discipline revises itself. Recent studies measuring AI’s applicability to historical work capture transmissive tasks while remaining blind to the interpretive core; approaches that identify textual markers of historical thinking detect symptoms that can be simulated, not the compositional process producing them. The consequences extend beyond the profession. When stochastic history circulates in classrooms, policy research and public media, non-specialists encounter pasts stripped of contingency and contestability – pasts that naturalize present arrangements rather than rendering them open to challenge. The defence of writing advanced here is methodological rather than nostalgic: It preserves the conditions under which historical claims can be scrutinized and revised.

In a recent *New Yorker* essay, historian of science D. Graham Burnett asked whether the humanities could survive the rise of artificial intelligence (AI).¹ He described a classroom experiment in which students, confronted with the ubiquity of systems such as ChatGPT, stopped asking whether AI use was ethical or practical. Their silence did not register disengagement. It indicated a more profound shift in epistemic expectation. If machines could simulate the outputs of critical thinking, what remained of the labour through which thinking occurred? More pointedly, what was the purpose of writing if the semblance of interpretive discourse could now be generated without the interpretive work?²

Burnett’s classroom intervention is significant because he incorporated AI as an object of critical analysis. Rather than banning generative systems, he institutionalized their presence as an epistemological obstacle that students had to confront. They were instructed to use large language models to generate essays on canonical theorists such as Walter Benjamin, Michel Foucault and Roland Barthes, and then to write critical commentaries diagnosing the conceptual failures embedded

[Correction added on 2 April 2026, after first online publication: Peer review history statement has been added.]

¹ D. Graham Burnett, ‘Will the humanities survive artificial intelligence?’, *The New Yorker*, 26 April 2025, <<https://www.newyorker.com/culture/the-weekend-essay/will-the-humanities-survive-artificial-intelligence>> [last accessed 9 February 2026].

² Special thanks to Professor Stephen Vlastos of the University of Iowa for prompting me to reflect on the implications of Burnett’s experiment and encouraging me to write about it.

in the AI-generated output. The exercise displaced authorship from originality to accountability. Students were not asked what they could say that machines could not. They were asked what they could see that the machine had missed. Diagnosing those failures required more than technical acumen. It required the capacity to articulate what counted as a valid interpretation, how such validity might be tested and what made a sentence meaningful within a disciplinary context.

Students who began with a chatbot-generated précis of Benjamin's 'The Work of Art in the Age of Its Technological Reproducibility' found that while key vocabulary such as aura, mechanical reproduction and ritual function appeared fluently, the dialectical tension that structures Benjamin's argument had been erased.³ One student remarked that the chatbot had flattened the dialectic, producing a version of Benjamin that was untroubled by contradiction and unable to sustain historical temporality. Another noted that the paraphrase of Barthes's 'The Death of the Author' treated authorship as a matter of legal ownership rather than as a textual function embedded in practices of reading and discourse.⁴ The students were not simply cataloguing errors. They were reconstructing the epistemic conditions under which claims make sense. What emerged was not a defence of literacy in a conventional sense but a demonstration that understanding cannot be simulated. It must be composed.

The urgency of this insight becomes clearer when set against recent attempts to quantify AI's occupational reach. In December 2025, Microsoft Research published an analysis of 200,000 anonymized conversations with Bing Copilot, mapping user interactions to the O*NET database (a US Department of Labor taxonomy that classifies occupations by their component tasks) and computing an 'AI applicability score' (a composite measure of how often and how successfully AI-assisted tasks associated with each profession) for each occupation.⁵ Historians ranked second among all occupations, with an applicability score of 0.462 – exceeded only by interpreters and translators. Writers, editors and journalists followed closely behind. The study concluded that AI proves most useful for 'information work', defined as 'the creation, processing, and communication of information', and that historians' work falls squarely within this domain.⁶

The finding appears alarming. It suggests that historical practice lies at the centre of what current AI systems can assist or perform. Yet the methodology reveals its own limitations. The O*NET framework decomposes occupations into discrete 'Intermediate Work Activities' (standardized task descriptions intended to allow comparison across professions) – tasks such as 'edit written materials', 'respond to customer inquiries', 'provide information to the public' and 'research historical or social issues'.⁷ When mapped to historians, the activities that drive the high applicability score are overwhelmingly *transmissive*: explaining, communicating, editing and maintaining knowledge. These are the activities that historians share with

³ Walter Benjamin, 'The work of art in the age of its technological reproducibility: Second version' (orig. 1935–36), in Michael W. Jennings, Brigid Doherty, and Thomas Y. Levin (eds), *The Work of Art in the Age of Its Technological Reproducibility, and Other Writings on Media* (Cambridge, MA, 2008), pp. 19–55.

⁴ Roland Barthes, 'The death of the author' (orig. 1968), in Roland Barthes, *Image Music Text: Essays Selected and Translated by Stephen Heath* (New York, 1977), pp. 142–8.

⁵ Kiran Tomlinson, Sonia Jaffe, Will Wang, Scott Counts, and Siddharth Suri, 'Working with AI: Measuring the applicability of generative AI to occupations', arXiv preprint arXiv:2507.07935v6 (22 December 2025), <<https://doi.org/10.48550/arXiv.2507.07935>>.

⁶ Tomlinson et al., 'Working with AI', p. 1. The authors define 'information work' as work involving 'the creation, processing, and communication of information'.

⁷ US Department of Labor, O*NET OnLine, <<https://www.onetonline.org/>> [accessed 10 February 2026]. The Intermediate Work Activities (IWAs) for historians include 'Editing written materials', 'Responding to customer/client inquiries or complaints', 'Providing information to the general public', and 'Researching historical or social issues'.

teachers, journalists and public information officers. They describe how historical knowledge circulates, not how it is constituted.

What the framework cannot capture is the recursive relationship between interpretation and composition that makes historical claims answerable. The study's own success metrics – user satisfaction and 'task completion' – measure whether Copilot produced outputs that users found adequate to their immediate purposes.⁸ A user who asks 'Why did the Tokugawa shogunate collapse?' and receives a fluent paragraph that explains familiar factors may register satisfaction. The paragraph may even be correct in the sense of reproducing established interpretations. What the metric cannot assess is whether the output carries marks of warranted explanation – whether it bears traces of sources consulted, periodizations weighed and causal paths tested. As the authors themselves acknowledge, the data suggest that 'AI can help users with a broader fraction of their work than it can perform directly'.⁹ The distinction matters. Assisting a historian who maintains interpretive control differs fundamentally from performing the interpretive work itself.

The Microsoft study thus inadvertently demonstrates what occupational frameworks systematically miss. Historians appear highly AI-applicable because their work involves activities – reading, writing, explaining, synthesising – that AI can simulate at the surface level. The study captures what historians *also* do: respond to inquiries, prepare informational materials, maintain expertise. It cannot capture what historians *distinctively* do: constitute knowledge through the friction of composition, render temporal structures legible, make methodological choices that remain traceable and contestable. The high applicability score reflects AI's capacity to occupy the transmissive perimeter of historical practice while leaving the interpretive core untouched – and unmeasured. This essay terms such output 'stochastic history': historical-sounding prose produced by probabilistic text systems that substitutes linguistic plausibility for the warrant of evidence. The term names not a deficiency of accuracy but a structural condition – prose that inherits the surface forms of historical explanation without enacting the inferential work through which claims become answerable.

Marnie Hughes-Warrington has recently proposed a complementary approach to this diagnostic problem. In her account, genuine historical writing exhibits what she calls 'healthy noise' – conditionals, counterfactuals, rhetorical questions and 'chronological flicker' that hold open alternative possibilities and signal openness to new evidence.¹⁰ AI-generated historical texts, she observes, tend towards confident assertion and lack this interrogative texture. Yet, as Hughes-Warrington herself acknowledges, AI systems may eventually learn to simulate these surface features. The question then becomes whether the markers of historical thinking can be reduced to textual patterns at all – or whether they depend on something that cannot be detected in the prose alone. This essay argues that because large language models sever historical prose from the methodological trail through which claims become answerable, composition itself – not merely its outputs – must be recognized as the practice that distinguishes history from its stochastic simulation.

⁸ Tomlinson et al., 'Working with AI', p. 3. The 'completion' metric uses an automated classifier (itself a large language model) validated against user feedback (thumbs up/down).

⁹ Tomlinson et al., 'Working with AI', p. 12.

¹⁰ Marnie Hughes-Warrington, 'Ethics for artificial historians', *History and Theory* 64, no. 2 (2025), pp. 159–77. Hughes-Warrington proposes scanning for 'historiographical logic' – a hybrid of modal, propositional and erotetic (question-based) reasoning – as a marker of ethically beneficent historical writing.

The stakes of stochastic history

The stakes of this confusion are not merely professional. When stochastic outputs circulate as adequate substitutes for historical explanation, claims lose their traceability—there is no methodological trail to audit, no position from which interpretive choices can be defended or contested. Temporal conflicts flatten into smooth chronology; the layered pasts and contested futures that historical thinking renders visible disappear behind sequences optimized for coherence. Dominant narratives, overrepresented in training data, recycle without the archival friction or peer contestation that enables revision. If the tools mediating education, journalism and policy research cannot distinguish between fluent simulation and warranted argument, public deliberation loses access to the histories that make present arrangements appear contingent rather than inevitable. The risk is not that AI will replace historians. It is that historical-sounding prose will displace the conditions under which historical claims remain inspectable.

In July 2025, the American Historical Association (AHA) approved Guiding Principles for Artificial Intelligence in History Education codifying this insight. The AHA's ad hoc committee began by affirming that generative AI cannot replace human teachers and warned against viewing automation as a remedy for the profoundly human work of education. Like Burnett's experiment, the AHA highlights the risks of hallucination and bias, reminding educators that large language models produce texts rather than truths. Rather than banning these tools – a solution the AHA deems impractical and unenforceable – it urges cultivating AI literacy.

In the United Kingdom, the Royal Historical Society (RHS) has run a programme on generative artificial intelligence that grew out of an acute recognition that historians should interrogate emergent technologies rather than either embrace them uncritically or reject them outright. From Adam Budd's pre-event essay for the RHS's July 2024 roundtable to a pilot study and the eventual release of a reading guide in October 2025, the RHS has repeatedly returned to questions of trust, authority and pedagogy: students found that AI answers lacked depth, context and evidence.¹¹ In September 2025, the RHS's workshop at the University of Edinburgh demonstrated its intent to transform these conversations into actionable guidance.¹²

Placed in dialogue, these transatlantic initiatives reveal both convergence and divergence. The RHS's reading guide and pilot study are pragmatic experiments that exposed the limits of generative AI and encouraged educators to treat AI as a starting point rather than an endpoint. The AHA document codified similar insights into broader pedagogical principles, emphasizing that banning AI is neither viable nor desirable and that educators must instead craft transparent policies and teach students how to cite AI outputs. The AHA went further by offering a model policy table delineating acceptable and unacceptable uses, from using AI for summarization to prohibiting the submission of AI-generated essays as one's own.¹³ In contrast, the RHS programme has yet to publish formal policies, but its Edinburgh workshop is expected to produce guidance notes.¹⁴ Together, these efforts signal a transatlantic movement within the historical profession: Both bodies recognize that generative

¹¹ Adam Budd, 'Gen AI, history and historians', Royal Historical Society blog, 11 July 2024, <<https://blog.royalhistsoc.org/2024/07/11/gen-ai-history-and-historians/>> [last accessed 9 February 2026].

¹² Royal Historical Society, 'Generative AI, history and historians, a reading guide', *Historical Transactions* (blog), first published May 2024, updated September 2025, <<https://blog.royalhistsoc.org/2025/10/02/generative-ai-history-and-historians-a-reading-guide/>> [last accessed 10 February 2025].

¹³ American Historical Association, 'Guiding principles'.

¹⁴ David Clayton, Henrice Altink, and Esther Wilson, 'Piloting responsible and effective use of generative AI in undergraduate history teaching', Royal Historical Society blog, 16 July 2025, <<https://blog.royalhistsoc.org/2025/>>

AI is reshaping the educational landscape and that responsible engagement requires critical experimentation, institutional guidance and reaffirmation of the discipline's foundational methods.

These parallel institutional responses point towards a deeper question. The central issue is not whether AI can write competent sentences but whether those sentences carry marks of warranted explanation that make historical writing legible as knowledge rather than as performance. This essay begins where these institutional frameworks end. Composition is treated here as an epistemic practice rather than as a channel for already formed ideas, and generative systems are analysed in terms of how they displace that practice. For this purpose, I return to the term 'stochastic history'. The sequences appear competent because they inherit patterns from prior discourse and recombine them into plausible order. They do not carry the marks of composition through which historians test causal paths, make periodization choices visible and acknowledge silences that shape interpretation. What results is text that reads like history yet lacks a methodological trail that a reader can audit. The sequence may persuade. The claim is not warranted.¹⁵

Burnett's experiment and the institutional principles help to locate AI as a classroom antagonist that clarifies the stakes of interpretation. His assignment shifts authorship from originality to accountability, asking students to diagnose machine failures. The AHA and RHS codify prudential norms for teaching with generative systems, warning that such tools hallucinate, urging disclosure and insisting that expertise must come first. This essay intervenes at a different register. It specifies what makes machine failure epistemic rather than merely rhetorical, and it identifies the locus of historical knowledge in the act of composition itself. The key distinction is the one formalized as plausibility versus warrant. Token prediction severs sentences from a methodological trail that a reader must be able to inspect. Historical prose, by contrast, must bear the traces of periodizations chosen and rejected, causal paths tested and abandoned and silences acknowledged as constitutive. Burnett's assignment organizes failure so that students can practice critique. The professional associations offer institutional guidance at the level of policy. My analysis explains why stochastic outputs are not history and proposes criteria for recognizing when writing bears the marks of historical labour. The displacement that this essay analyses is also the rupture from which it writes: Stochastic history makes visible the relationship between composition and knowledge that the discipline had not previously needed to defend against external dissolution – because no technology had been capable of severing the sentence from the labour that produced it.¹⁶

Machines, sequences and the absence of interpretation

The epistemological problem begins with how large language models currently operate. Systems such as OpenAI's ChatGPT, Anthropic's Claude, Google's Gemini and the increasingly popular Perplexity AI – all built on transformer architectures trained via next-token prediction (i.e., generating each word by calculating the most statistically likely continuation) – do not reason, select or interpret in any disciplinary

07/16/piloting-responsible-and-effective-use-of-generative-ai-in-undergraduate-history-teaching/> [last accessed 9 February 2026].

¹⁵ Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Margaret Mitchell, 'On the dangers of stochastic parrots: can language models be too big?', in *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency* (2021), pp. 610–23, <<https://doi.org/10.1145/3442188.3445922>>; Hayden White, *Metahistory: The Historical Imagination in Nineteenth-Century Europe* (Baltimore, MD, 1973).

¹⁶ Burnett, 'Will the humanities survive artificial intelligence?'; American Historical Association, 'Guiding principles'; Clayton, Altink, and Wilson, 'Piloting responsible and effective use'.

sense. They generate each word by maximizing conditional probability based on exposure to vast corpora of written language. This allows them to produce sentences that conform syntactically to genre expectations and even to emulate the discursive rhythm of historical prose.¹⁷ When prompted to explain the collapse of the Tokugawa *bakufu*, the military government that ruled Japan from 1603 to 1868, a stochastic system might compose a fluent paragraph that references rural unrest, fiscal strain, samurai dissatisfaction and foreign pressure.¹⁸ The sequencing will appear historical. The vocabulary will align with recognizable tropes. The transitions will flow. None of the research, interpretation or methodological deliberation that makes such prose accountable will have taken place.

The system will not have consulted sources, traced historiographical debates, weighed rival periodizations or acknowledged gaps in the evidence. It will not have deliberated over how to render causality visible, which silences require marking, or how to determine the temporal architecture of the collapse. What appears instead is what Bender and her coauthors describe as grammatical plausibility without genuine meaning: text that reads as meaningful because it inherits patterns of coherence but carries no interpretive burden.¹⁹ A historian's sentence, such as 'rural protest intensified as grain prices rose', implies an entire chain of operations. Protest events have been identified and dated. Price data have been collected and aligned with those dates. Rival explanations have been considered. A causal logic has been articulated and tested against the archive.

Not every historian performs all these operations for every sentence. Historians working under deadline pressure, reproducing seminar templates or following supervisors' framings without critical engagement, may generate similarly unreflective prose. The difference does not lie in the string of words. It lies in the relationship between the text and a recoverable process. The student's work, however thin, remains answerable. An MA dissertation or PhD thesis supervisor can ask which sources were consulted, why a particular periodization was chosen and what alternative explanations were considered. The student may answer poorly, revealing shallow engagement, but there remains a person who can be asked. With machine-generated prose, no such questioning is possible. The text has no author in the minimal sense required for accountability. There is no position from which interpretive decisions could have been made and might now be defended.

This boundary matters pedagogically and epistemologically. An underprepared student occupies a position within a system of accountability that can detect and correct unreflective work. Stochastic history, by contrast, circulates without institutional friction. It appears competent precisely because it inherits patterns from competent writing, which makes detection harder, not easier. The pedagogical challenge intensifies. Instructors must teach students to recognize when apparent competence masks absent accountability, including in their own work. The relevant test is not textual polish but structural answerability. One can diagnose stochastic history by asking which sources support a claim, how competing periodizations were weighed and what causal paths were tested and rejected. When these questions produce empty answers, when no methodological position exists from which they could be answered, the prose is stochastic regardless of its surface coherence.

¹⁷ Bender et al., 'On the dangers of stochastic parrots'.

¹⁸ Robert Hellyer, 'Foreign Relations and Coastal Defense under the Mature Tokugawa Regime', in David L. Howell (ed.), *The New Cambridge History of Japan*, vol. 2: *Early Modern Japan in Asia and the World, c. 1580–1877* (Cambridge, 2024), 159–83.

¹⁹ Bender et al., 'On the dangers of stochastic parrots'.

These limitations describe current systems, which generate text by predicting the most statistically likely next word based on fixed bodies of training data. Whether future systems that combine source databases, iterative revision loops and provenance tracking could approximate some operations of historical thinking remains an empirical question. Such architectures would differ fundamentally from the language models analysed here. They would require real-time access to sources, mechanisms for detecting evidential insufficiency and the capacity to revise under pressure rather than simply regenerate alternatives. Whether such modifications would constitute historical thinking or only sophisticated simulation raises philosophical questions beyond the scope of this essay. What matters for the argument is that current systems lack these capabilities and that even if future systems were to develop them, the pedagogical imperative would remain unchanged. Students must learn historical thinking through composition, not through tool use, however sophisticated the tools.

Recent empirical studies confirm this pattern. Björn Kindenberg's comparison of ChatGPT-generated and student-written narratives for lower-secondary history prompts found that machine-generated essays displayed surface coherence but lacked analytical structure, particularly in their treatment of causation and in the interpretation of sources.²⁰ Another study on undergraduate writing behaviour noted that students who used AI frequently prompted systems to write the whole essay, bypassing not only the mechanics of composition but also the cognitive work of argument formation and evidence evaluation.²¹ The York pilot study found similar patterns: Students who engaged with Google Gemini recognized that its outputs lacked the depth and evidential grounding required for historical argument.²² The challenge intensified in the mid-2020s as AI began to mediate research itself. By 2025, tools such as Perplexity and Google's AI Overview generated synthetic summaries in response to search queries, so machine interpretation now interposes itself between searcher and sources before composition even begins. Even where such systems retrieve and cite sources, the generative mechanism that assembles those citations into prose remains predictive rather than interpretive: selection is governed by relevance ranking, not by the evidential friction through which historians are trained to test claims.

A thought experiment clarifies what is at stake. A prompt such as 'Explain the causes of the Rice Riots of 1918 in Tokyo and Osaka' might yield a fluent answer that references rice inflation, urban protest and the state's response.²³ A meaningful test would ask the system to defend those claims. What sources support the price data? How were protest actors identified? Were regional variations considered? Which causal paths were tested and rejected? A human historian, faced with these questions, would gesture to documents, interpretive decisions and methodological frameworks. A generative system can only regenerate, offering variant phrasings of the same statistical association. There is no evidentiary trail to audit, no contested terrain through which the interpretation was forged. What is lost is not merely accuracy, but answerability.

The absence of a methodological position, of someone who made choices that can be interrogated, renders the prose orphaned. It is coherent in form but disconnected from the interpretive process. Readers accustomed to historical writing

²⁰ Björn Kindenberg, 'ChatGPT-generated and student-written historical narratives: a comparative analysis', *Education Sciences*, 14/5 (2024), 530.

²¹ Andrew Jelson, Daniel Manesh, Alice Jang, Daniel Dunlap, and Sang Won Lee, 'An empirical study to understand how students use ChatGPT for writing essays', version 2, preprint, 26 May 2025, arXiv, <<https://doi.org/10.48550/arXiv.2501.10551>>.

²² Clayton, Altink and Wilson, 'Piloting responsible and effective use'.

²³ Michael Lewis, *Rioters and Citizens: Mass Protest in Imperial Japan* (Berkeley, CA, 1990).

bring interpretive generosity to what they read. They assume that a coherent paragraph reflects intentionality, deliberation and accountability. With machine-generated prose, generosity becomes a liability. The problem intensifies beyond specialist audiences. When stochastic history circulates in policy briefs, journalistic summaries or educational materials, readers lack the disciplinary training to recognize absent accountability. A policymaker encountering an AI-generated paragraph about postwar land reform cannot distinguish statistical plausibility from evidence-based argument. Neither can the undergraduate encountering it in a course reading, nor the journalist synthesizing it for a wider audience. The specialist reader can learn to audit methodological absence. The non-specialist inherits conclusions without access to the interpretive labour that would make them trustworthy. This creates new stakes for public history. The problem is not only factual accuracy, but the visibility of how historical claims are made. As Bender and her coauthors observe, the coherence of language-model output misleads users into attributing meaning and intentionality to processes that possess neither.²⁴

Why it matters if machines think historically

When historians speak of historical thinking, they name a set of practices that link explanation to evidence through choices about time, narrative form and responsibility. Non-historians encounter the same issues in a different register. A student who asks a chatbot why the Japanese empire collapsed, a policy analyst who requests lessons from Weimar for contemporary democratic backsliding or a journalist who uses an AI assistant to draft a story on decolonization all receive more than isolated facts. They receive miniature narratives that select beginnings and endings, attribute causality and connect past events to present concerns. These operations resemble historical explanation, yet they are governed by statistical pattern-matching across a fixed body of training text rather than by evidentiary labour and methodological debate.²⁵ Even if such systems cannot think in a human sense, they already shape how millions of people imagine the past. The question of whether machines think historically is therefore not peripheral. It concerns who sets the terms on which temporal reasoning will be automated and scaled and whether that automation will preserve or erase the conditions under which historical claims remain answerable.

For non-specialists, the first consequence of stochastic history is the distribution of responsibility and causation. Historical thinking trains readers to situate outcomes within chains of decision, structure and contingency so that present inequalities and conflicts appear as the result of identifiable processes rather than as natural conditions. Reconstructing the histories of industrial capitalism, colonial extraction or occupation governance is part of an effort to show how specific arrangements of power arose and how they might be challenged or undone. Stochastic systems that treat the past as a flat database of interchangeable statements tend instead to produce context-free generalities. Structural histories are compressed into moralizing commonplaces about unstable regions, authoritarian cultures or ancient hatreds. This flattening is not politically neutral. It blurs distinctions between perpetrators and victims, obscures long-running patterns of dispossession and naturalizes existing distributions of power by presenting them as the inevitable outcome of diffuse tendencies rather than as the result of contestable decisions. When tools that mediate policy research, journalism and education default to such dehistoricized framings, they quietly amplify those distortions even when individual users seek nuance.

²⁴ Bender et al., 'On the dangers of stochastic parrots'.

²⁵ Ibid.

Temporality forms a second site of risk and connects directly to the analytic threads developed later in the essay around anachronism, smoothing and the frozen archive. Reinhart Koselleck's distinction between the space of experience and the horizon of expectation highlights the layered, asymmetrical character of historical time, in which actors draw on inherited pasts while imagining divergent futures.²⁶ François Hartog's notion of regimes of historicity extends this insight by showing how communities inhabit different structures of temporal reference that shape what counts as agency, continuity or rupture.²⁷ Historians labour to reconstruct these temporal structures and to make visible how periodizations, anticipations and memories shape action. Current language models, by contrast, learn time as adjacency in text. They register which dates cluster with which labels and events, but they have no internal requirement to distinguish analytic period concepts from loose metaphor nor to notice when a modern category, such as populism or authoritarianism, is projected backwards. If such systems are allowed to stand in for historical explanation without being engineered to mark these differences, they will normalize presentist drift and routine anachronistic analogy. Non-historians will encounter pasts already assimilated to contemporary vocabularies, with little sense that period labels and temporal cuts are themselves contested.

A third issue concerns trust and disclosure. One of the discipline's core contributions to public life lies in its techniques for handling uncertainty. Historians use footnotes to distinguish evidence from inference, to discuss archival gaps that limit what can be claimed and to acknowledge disagreement among specialists. These practices do not guarantee correctness, but they make visible the conditions under which claims are advanced and invite readers to contest them. Stochastic systems, unless deliberately constrained, do the opposite. They smooth over uncertainty, present composite views as if they were stable consensus and rarely disclose when a claim rests on sparse, skewed or conflicting data.²⁸ A machine that thought historically in even a minimal sense would need to mark its own limits, indicate which sources it draws upon and expose users to contending interpretations rather than averaging them away. Without such capacities, widespread reliance on AI writing assistants risks training users to confuse fluency with warrant and to accept neatly packaged accounts in place of contested explanation. The concept of stochastic history names the endpoint of this trajectory: sequences that look historical while disabling the friction that makes claims auditable.

Hughes-Warrington's concept of 'healthy noise' illuminates one dimension of this problem. Prize-winning historians, she shows, deploy conditionals, counterfactuals and rhetorical questions at rates that AI-generated texts do not match. These features signal what she calls 'openness toward the possibilities of new evidence from the past, new ways of interpreting that evidence'.²⁹ Yet openness is itself a posture that can be performed. A stochastic system prompted to hedge its claims, introduce counterfactuals and pose rhetorical questions would produce text exhibiting Hughes-Warrington's markers without the interpretive friction those markers are meant to index. The conditional 'perhaps Toussaint's grandfather was a provincial governor' carries different epistemic weight when it emerges from an archival gap the historian has actually encountered than when it is generated because conditionals pattern well in historical prose. What distinguishes warranted from simulated uncertainty is not

²⁶ Reinhart Koselleck, *Futures Past: On the Semantics of Historical Time*, trans. Keith Tribe (New York, 2004).

²⁷ François Hartog, *Regimes of Historicity: Presentism and Experiences of Time*, trans. Saskia Brown (New York, 2015).

²⁸ Bender et al., 'On the dangers of stochastic parrots'.

²⁹ Hughes-Warrington, 'Ethics for Artificial Historians', 169.

the presence of hedging language but the compositional process through which that language came to be selected – a process that leaves traces in the architecture of the argument, not merely in its word choice. Hughes-Warrington's diagnostic is valuable precisely because it identifies *symptoms* of historical thinking. But symptoms can be mimicked. The *condition* that produces them – the recursive testing of interpretation against evidence during composition – cannot.

These concerns feed into a broader democratic stake. Public deliberation about climate reparations, migration regimes, memorialization or technological regulation depends on citizens who can place proposals within longer histories of policy, power and struggle. If the tools that structure everyday reading and writing treat history as a reservoir of illustrative anecdotes rather than as an argumentative discipline, they will narrow the imaginative horizon of what people take to be possible or sayable. In that setting, the risk is not only that AI will misdescribe particular episodes. It is that it will habituate readers to a style of historical speech in which conflict, contingency and alternative futures disappear behind a smooth surface of lessons from the past. Teaching machines to think historically is not a naïve claim that code can acquire consciousness. It names an attempt to encode, as far as current techniques permit, some of the profession's hard-won habits of attention to provenance, sensitivity to temporal scale, explicit handling of uncertainty and acknowledgement of conflict over interpretation.

The institutional principles of the AHA and RHS gesture in this direction by insisting that generative systems produce texts rather than truths and that historical expertise must precede reliance on such tools.³⁰ The argument here specifies why. Without the compositional labour through which historical claims become answerable, these systems cannot sustain the discipline's commitments however polished their sentences may appear. To ask why it matters if machines think historically is therefore to ask how far historians are willing to let probabilistic systems determine the forms in which the past becomes thinkable in classrooms, policy offices and public culture.

Historical time and the collapse of temporal judgement

If stochastic history evacuates interpretive labour, the first casualty is temporal judgement, the ability to read and render the layered, competing temporalities that shape historical events. Koselleck's work on historical temporality offers an analytic anchor for understanding what disappears when stochastic systems generate prose. His distinction between the space of experience and the horizon of expectation foregrounds the layered and asymmetrical nature of historical time.³¹ Events do not unfold within a single temporal stream. They emerge from intersections of inherited pasts and imagined futures. A historian writing about postwar labour activism in Japan must attend to multiple temporalities simultaneously. The *longue durée* of industrial capitalism, the immediate aftermath of wartime mobilization, legal ruptures introduced under Allied occupation and the anticipated futures projected by labour organizers, bureaucrats and corporate actors all bear on what counts as agency and how periodization is configured. A labour organizer who sees occupation reforms as breaking wartime coercion inhabits a different temporal regime than a bureaucrat

³⁰ American Historical Association, 'Guiding principles'; Clayton, Altink, and Wilson, 'Piloting responsible and effective use'.

³¹ Reinhart Koselleck, "'Space of experience" and "horizon of expectation": Two historical categories', in Koselleck, *Futures Past*, 255–75.

who fears that those same reforms repeat prewar instability. The historian must render both temporalities legible.

Hartog extends this analysis with his concept of regimes of historicity, socially embedded frameworks that structure how communities relate pasts and futures.³² Unlike the chronological lines presented in textbooks, such regimes do not merely order events. They construct differential relationships to time itself. An activist for whom the past figures as an unfinished struggle occupies a different regime than a bureaucrat who imagines a developmental arc towards modernization. These distinctions matter because, at its best, historical writing renders them visible. Take, for example, a sentence such as 'While Supreme Commander for the Allied Powers (SCAP) reforms promised democratic restructuring, many rural labour groups feared a repetition of prewar co-optation'. The statement in question enacts temporal multiplicity. One clause captures reform as a future-oriented break. The other captures reform as a continuation of past domination. The sentence is not merely informative. It stages temporal conflict.

Current stochastic systems do not perform this work. They process sequences based on statistical likelihood rather than temporal structure. Their architectures provide no mechanism for distinguishing chronological adjacency from temporal resonance or for marking the simultaneity of non-synchronous experiences that Koselleck regarded as constitutive of modernity.³³ When prompted to describe postwar labour conflict, a model might generate sentences that appear to flow chronologically, inserting dates, events and actors in familiar arrangements. What underpins that fluency is not temporal judgement but token prediction. Returning to the Japanese example, the model cannot account for the possibility that rural workers may have interpreted land reform through memories of tenancy disputes from the 1920s, while American policymakers understood reform as a Cold War imperative. It cannot capture the tension between inherited experience and emergent expectation.

The result is anachronism by smoothing. Temporal conflicts are levelled into adjacency. One can diagnose this failure by asking which workers expected democratization, under what conditions, drawing on which pasts and anticipating which futures. When such questions produce empty answers, when expectation floats unmoored from the regimes of historicity that would make it intelligible, temporal complexity has been evacuated. The system reproduces vocabulary and sequence but loses the conflicts that generate historical meaning. A generated sentence, such as 'workers expected democratization' may be grammatically correct, but it lacks situational anchoring. The prose simulates continuity. The historical time it gestures towards has been collapsed.

Walter Benjamin warned against this kind of continuity. In his thesis on the philosophy of history, he described the historian's task as seizing ruptures rather than smoothing the flow of time, interrupting the apparent continuum to expose the stakes of memory and power.³⁴ Generative systems, as presently designed, do the opposite. They cannot interrupt inherited temporal logics. They cannot write against dominant periodizations. They reproduce frequency patterns from their training data, inherit continuity as default and recombine sequences into legible flows. The discipline's revisionary engagement with time, visible in disputes over when the postwar begins or how to periodize empire and decolonization, has no correlate

³² Hartog, *Regimes of Historicity*, 1–19.

³³ Reinhart Koselleck, 'History, histories, and formal time structures' in Koselleck, *Futures Past*, pp. 93–104 as well as idem, 'Space of experience' and 'horizon of expectation'.

³⁴ Walter Benjamin, 'Theses on the philosophy of history', in Hannah Arendt (ed.), *Illuminations*, transl. Harry Zohn (New York, 1968), pp. 253–64.

in systems that treat dates as tokens and temporal structures as side effects. Yet the arrival of these systems is itself a rupture that the discipline can seize – one that makes visible the work of composition precisely by threatening to render it unnecessary.

Narrative form as epistemic commitment

Hayden White's account of narrative emplotment clarifies how form and explanation are bound together. The narrative mode through which historical events are rendered, whether tragic, romantic, satirical or comic, shapes not only readers' expectations but the attribution of agency and causality.³⁵ To emplot the end of the Tokugawa shogunate as tragedy is to suggest inevitability and loss, to emphasize structural contradictions that made collapse unavoidable. To cast it as romance is to frame it as redemption and renewal, to locate agency in reformist actors who overcame an exhausted regime. The formal choice influences which evidence is collected and what counts as an explanation. Narrative decisions are not decorative. They are interpretive commitments with significant consequences.

Language models do not choose emplotment in this sense. They inherit it. They lack an internal mechanism for deciding whether a story should resist closure or whether sources demand impasse rather than resolution. They optimize for coherence within the distribution of forms present in their training data. If those data disproportionately represent reconciliatory narratives, as is plausible given the structure of encyclopaedias, textbooks and many syntheses, generated prose will tend towards closure even when the archive resists it. A model of the dissolution of the Japanese empire is likely to end with recovery, democratization or an economic miracle, even if primary documents reveal protracted repatriation crises, unresolved violence in former colonies or persistent structural exclusion of colonial subjects returned to the metropole. The system does not know how to ask which narrative form the materials require. It selects what seems likely.

This is not only a matter of bias. It is a failure of judgement. Hartog reminds us that to write history is to position oneself within a temporal regime that renders past, present and future intelligible in a particular way.³⁶ That position entails decisions about what to emphasize, what to withhold and how to frame stakes. Historians internalize narrative forms through training, but the discipline expects them to reflect on those forms and revise them under pressure. A graduate student may instinctively reach for an arc that moves from protest to repression to compromise. In confronting sources that resist that arc, evidence of unresolved militancy or incoherent outcomes, the student is forced to adjust. If the archive reveals fragmentation rather than settlement, the historian cannot conclude that compromise was reached without falsifying the evidence. Form must respond to structure. Revision is not aesthetic. It is epistemic.

Language models regenerate rather than revise. They vary their output. They offer alternative phrasings. They do not revise under pressure from sources. This describes current architectures that lack feedback loops between outputs and evidential material. Whether future systems might implement such loops does not alter the present claim. Generated history lacks accountability. Models do not return to archives. They do not reperiodize. They do not halt when a claim cannot be sustained. Readers encountering such prose may recognize familiar emplotment. The familiarity

³⁵ Hayden White, *Metahistory: The Historical Imagination in Nineteenth-Century Europe* (Baltimore, MD, 1973); Hayden White, *The Content of the Form: Narrative Discourse and Historical Representation* (Baltimore, MD, 1987).

³⁶ Hartog, *Regimes of Historicity*, pp. 1–19.

conceals a more profound absence, the absence of decision, of judgement, of interpretive struggle.

Writing as method: Composition and the epistemic function of labour

From divergent methodological positions, E. H. Carr, G. R. Elton and Marc Bloch all insisted that writing is not merely the medium through which interpretation is conveyed but the site where interpretation is constructed.³⁷ Carr argued that facts do not speak for themselves but require selection and framing. Elton claimed that the argument emerges through the disciplined exposition of sources. Bloch described the historian as an ogre who follows the scent of human flesh, meaning that history is an inquiry into human experience that becomes knowledge only when staged in prose that makes the process visible. Writing is where facts become evidence, where hypotheses are tested and either discarded or transformed into an explanation.

Lynn Hunt makes a related point in her reflection on writing as thinking. Historians may believe they know what they mean, but it is only in the process of drafting, revising and confronting contradictions on the page that their claims acquire clarity and defensibility.³⁸ An attentive historian who writes ‘that policy failed because elites resisted’, but cannot specify which elites, what forms of resistance or why the opposition mattered, discovers through writing that the claim requires repair. That repair is not stylistic. It is cognitive. Peer review, advisor feedback and self-critique reveal where prose has outrun evidence, where an interpretation cannot be sustained or where a narrative arc has been imposed rather than earned.

The paragraph that will not cohere, where transitions fail, evidence contradicts or temporal logic breaks becomes diagnostic. It shows where explanation was presumed rather than developed from sources. The historian must then return to archives, reconsider periodization or acknowledge that the archive will not support the claim. Composition functions as a proving ground. It is where claims meet constraints, where interpretive ambition collides with archival refusal and where the historian discovers whether an explanation holds under pressure. Writing in this sense is not a report of understanding formed elsewhere. It is the medium in which understanding is forged.

Stochastic systems assemble rather than compose. During training, they learn probability distributions from text corpora. During generation, they predict the next likely token. They reproduce genres of historical prose because such patterns are statistically frequent, but they do not engage in recursive deliberation that transforms evidence into argument. A system may state that ‘urban protest intensified in response to food shortages’ without asking whether this claim withstands conflicting reports, whether contemporaries understood causality similarly or whether the evidence points instead to labour organization or political repression. The difference does not lie in the presence or absence of text. It lies in process. The historian writes into uncertainty. The machine writes through it.

Pedagogies of friction: Burnett, Martinsen and the architecture of judgement

The pedagogical models developed by Burnett and Martinsen offer strategies for preserving writing as a site of interpretation in an educational landscape that

³⁷ Marc Bloch, *The Historian's Craft*, transl. Peter Putnam (Manchester, 1954); E. H. Carr, *What Is History?* (London, 1961); G. R. Elton, *The Practice of History* (London, 1967).

³⁸ Lynn Hunt, ‘How writing leads to thinking’, *Perspectives on History*, 1 February 2010, <<https://www.historians.org/perspectives-article/how-writing-leads-to-thinking-february-2010/>> [last accessed 9 February 2026].

now includes generative systems as default tools. Burnett's assignment, which requires students to critique AI-generated essays, foregrounds failure as an epistemic stimulus.³⁹ Students do not merely list mistakes. They reconstruct the conditions under which interpretation fails. When a chatbot flattens Benjamin's dialectics into a chronology of technological reproduction, students must articulate what is missing: tension, rupture and the historical specificity of aura as a mode of cultural memory under capitalism. Their commentaries become acts of recovery, not of content, but of method.

A complication follows. Students could, in principle, use AI to generate their critical commentaries, asking another system to critique the first. If the assignment succeeds by staging friction between machine fluency and human judgement, that friction dissolves when judgement itself is delegated. Martinsen anticipates this problem by requiring students to document their process, including drafts, revisions and moments of confusion, so that the labour becomes visible rather than only the final product.⁴⁰ Assessment shifts towards the trace of thinking. Instructors must examine not only argumentative quality but also signs of compositional struggle. This is labour-intensive and difficult to scale, which helps explain why such assignments remain experimental.

Martinsen presses further by designing tasks that scaffold reflection into each stage of writing. Students begin with a freewrite, then engage with an instructor-tuned chatbot that surfaces counterclaims or questions assumptions. Students annotate its responses, revise their arguments and submit a final piece alongside commentary that documents their cognitive arc. Here, revision is not a mark of failure. It is evidence. The final essay is judged not only on polish but on the interpretive path that produced it. In this model, writing becomes inspectable as a process. It records how interpretation was tested, reconsidered and refined. Composition becomes curriculum rather than a mere instrument.

Both models refuse to separate composition from cognition. Writing is not the endpoint of thought. It is the terrain on which thinking occurs. Martinsen remarks that institutions claimed to value process but rewarded polish and that AI exposes the gap between those claims and assessment practices.⁴¹ His response is to realign pedagogy with the epistemic stakes of the discipline. If machines can simulate fluency, instruction must return to friction, to the struggle for coherence when evidence resists, to the discomfort of discarding an initial interpretation, to the exposure that occurs when a claim cannot be sustained.

These experiments accept AI's presence in students' lives while rejecting the premise that writing is reducible to text production. Writing becomes the record of interpretive effort, the means by which students learn to bear the weight of historical claims. They also indicate how fragile that record becomes when generative systems are treated as harmless drafting aids. If students increasingly outsource not only sentences but argumentative structure to stochastic systems, the visible traces of interpretation begin to disappear. Pedagogies of friction attempt to restore those traces, but they do so within institutions that often lack the resources to support such labour at scale.

³⁹ Burnett, 'Will the humanities survive artificial intelligence?'

⁴⁰ Eric Lars Martinsen publishes 'Notes on teaching writing in the AI era', on Substack, where he regularly comments on these issues: <https://elmartinsen.substack.com>. In particular, see Eric Lars Martinsen, 'Vigor, Voice, and the Future of College Writing in the Age of A', 14 April 2025, <https://elmartinsen.substack.com/p/vigor-voice-and-the-future-of-college>; and Eric Lars Martinsen, 'Process over product in college writing: A human-centred approach to writing instruction with AI', <https://elmartinsen.substack.com/p/process-over-product-in-college-writing>. [both accessed 9 February 2026].

⁴¹ Martinsen, 'Process over product in college writing'.

Authorship, answerability and the disappearance of method

Poststructuralist critiques by Roland Barthes and Michel Foucault dismantled the notion of the sovereign author as the origin of meaning. Barthes's famous declaration that the birth of the reader must be paid for with the death of the author and Foucault's analysis of the author function as a mode of discursive regulation have often been read as undermining appeals to authorial intention.⁴² Historians can accept these critiques without abandoning accountability. Foucault's author function describes institutional attribution and disciplinary regulation, not the elimination of responsibility for how evidence is used. Barthes shifts attention from biographical intent to textual effect, not from explanation to indifference.

The point here is that historical claims must remain traceable to decisions about evidence, interpretation and judgement, whether those decisions are associated with an individual name or with a collective practice. What matters is not the presence of an author as a psychological subject but the availability of a method. A historical claim about Meiji education, wartime mobilization or postwar land reform must be traceable to decisions about what counts as a source, which voices to prioritize, how to periodize and how to articulate causality. The author need not be sovereign. The argument must be answerable.

Generative systems eliminate this possibility. They produce sentences that simulate historical assertion without any position from which those assertions can be defended. There is no methodological agent, no research trail, no revision history. When a reader encounters such prose and attempts to reconstruct the logic that produced it, they find only the illusion of interpretive process, the assumption that coherence reflects labour. The explanation evaporates. Meaning becomes a projection. Readers attribute the method to what is, in fact, the result of statistical optimization. This is not the death of authorship in Barthes's sense. It is the disappearance of method. With it goes the possibility of historical speech in the disciplinary sense, in which claims can be scrutinized, challenged and revised.

The frozen archive: Hegemonic framing and the absence of reflexivity

Emily Bender and her collaborators have shown that language models trained on massive internet corpora reproduce dominant perspectives as a structural effect of their design.⁴³ The digital archive is stratified by access, participation and platform regulation. Voices from historically marginalized communities are underrepresented. Alternative framings are filtered out as unreliable or toxic. What remains is a normalized dataset encoded with hegemonic assumptions about identity, agency and causality.

When models learn from such data, they inherit a worldview. A system that asked about the United States occupation of Japan, trained on Wikipedia, State Department histories and widely used textbooks, is likely to emphasize reconstruction, democratization and constitutional reform. It is less likely to foreground sexual violence by occupation forces, racial hierarchy in military governance or the suppression of radical labour movements, such as the planned general strike of February 1947, which the occupation authorities banned.⁴⁴ These

⁴² Barthes, 'The death of the author'; Michel Foucault, 'What is an author', in *Language, Counter-Memory, Practice: Selected Essays and Interviews*, ed. Donald F. Bouchard, trans. Donald F. Bouchard and Sherry Simon (Ithaca, NY, 1977), pp. 113–38.

⁴³ Bender et al., 'On the dangers of stochastic parrots'.

⁴⁴ For prominent historiographical treatments of these themes, see John W. Dower, *Embracing Defeat: Japan in the Wake of World War II* (New York, 1999); Sarah Kovner, *Occupying Power: Sex Workers and Servicemen in Postwar*

omissions are not malfunctions. They reflect training data composition. The system optimizes for coherence within a skewed corpus.

Historians also reproduce hegemonic framings, often unconsciously and sometimes under constraints. A Japanese historian writing about wartime mobilization may inherit narratives of national sacrifice without examining gendered coercion in factory labour. An American historian of the occupation may emphasize democratization while eliding racial hierarchies within the military government. What distinguishes historical writing is the existence of mechanisms that make resistance possible. Three such mechanisms matter for present purposes, even though they operate unevenly across the profession.

The first is archival friction. The historian at a well-funded institution, with language training and extended fieldwork, encounters documents that contradict received narratives. The contingent faculty member without funding or access may reproduce textbook frameworks not out of complacency but out of constraint. In either case, an archival encounter rarely confirms a hypothesis in straightforward ways. Sources yield contradictions, silences and recalcitrant detail. A historian writing about the 1960 Miike coal strike may begin with a structural framework of labour–state conflict and discover internal factionalism, gendered exclusions or local memories that unsettle neat causal arcs.⁴⁵ Such encounters force revision. Transitions stumble when the prose attempts to move from unified class action to documented internal conflict. Paragraphs collapse when causal claims meet evidence of competing motivations. Composition brings these frictions to light.

Language models universalize the condition of frozen access. They are trained on a fixed corpus and cannot encounter new documents during generation. They cannot discover contradiction because they optimize for coherence. Asked about Miike, a model might reproduce a standard structural account of militant labour and state repression, but it cannot stumble over the letters, police reports or oral histories that complicate that narrative. Those materials may be present in the training data as tokens, but they are not encountered as resistance during composition because there is no composition.

The second mechanism is peer contestation. Historians who publish in refereed venues submit their work to review that challenges framing and demands evidential specificity. Reviewers contest periodization, point to missing archives or propose alternative causal paths. Authors revise in response, return to sources and refine arguments. Historians writing textbooks, policy briefs or public essays encounter weaker review or none at all, yet even there, professional norms and informal critique exert some pressure. Language models do not experience such pressure at the point of generation. They are trained once, updated periodically and then deployed at scale. Their outputs are not revised in light of criticism of particular claims. When an AI-generated paragraph about postwar policy contains a questionable periodization or causal explanation, nothing compels the system to adjust.

The third mechanism is historiographical consciousness. Senior scholars write with awareness of how interpretive frames have been contested and revised. They remember when social history displaced political history, when cultural history challenged structuralism and when postcolonial critique exposed Eurocentrism. Graduate students and early-career historians are still learning these debates. They may reproduce dominant narratives without recognizing their genealogy, but they

Japan (Stanford, CA, 2012); Christopher Gerteis, *Gender Struggles: Wage-Earning Women and Male-Dominated Unions in Postwar Japan* (Cambridge, MA, 2009).

⁴⁵ Andrew Gordon, *The Wages of Affluence: Labor and Management in Postwar Japan* (Cambridge, MA, 1998).

can be taught. Reading Yoshimi Yoshiaki on comfort stations disrupts state-centric accounts of imperial expansion.⁴⁶ Komagome Takeshi's work on Japanese colonial education in Korea reveals how the rhetoric of 'civilising' obscured the political and cultural domination inherent in the lived experience of colonization.⁴⁷ Such encounters alter framing. They show that every narrative choice carries historiographical weight.

Language models inherit frequency distributions that encode past contests as resolved consensus. They regenerate that consensus as if no alternative had existed. Asked about Japanese imperialism, a system trained on English-language encyclopaedias and textbooks is likely to emphasize modernization and development while muting violence and extraction. The system reproduces a hegemonic narrative not because it evaluated competing frameworks and selected one but because that narrative appears most often in its corpus. Human reproduction of hegemony can be interrupted by education, critique and archival encounter. Algorithmic reproduction is structurally tied to the composition of training data and to predictive mechanisms that lack the capacity for reflexive revision.

Generative AI, therefore, risks freezing historical understanding. Without the resistance of the archive, systems cannot discover what complicates received narratives. Without peer contestation, they cannot be forced to refine their framings. Without historiographical consciousness, they cannot recognize that their syntheses repeat what was already dominant. The result is not a neutral summary but an active retrenchment. Models recycle inherited explanations and present them as if they were current knowledge. The danger is not that these systems replace historians. It is that their outputs begin to circulate as sufficient substitutes for historical thinking, displacing the mechanisms through which the discipline revises itself.

The irreducible work of historical writing

What must be preserved is writing as infrastructure. A historian who writes 'land reform as destabilizing rural hierarchies' has not simply reported a fact. She has consulted land registers that show property transfers, traced these across prefectures to establish patterns, weighed contemporaneous testimonies from landlords and tenants against retrospective accounts that reframe events through postwar democracy narratives, considered rival explanations that attribute change to urbanization or market forces and constructed a sentence that renders that interpretive labour open to scrutiny. The sentence carries within it the possibility of challenge. Another historian may contest the periodization, question the sources or propose an alternative causal path. The sentence's form permits these challenges. It names actors, provides evidence and suggests testable causality. Contestability is not incidental. It is what makes the claim historical rather than merely plausible.

The defence of writing must therefore be methodological rather than nostalgic. To insist that historians write is not to romanticize craft or the solitary writer. It is to preserve the conditions under which historical claims remain traceable to evidence and open to challenge. The operations that make historical explanation open to inspection and revision currently require compositional friction that stochastic systems do not produce. Whether future architectures might transform this situation

⁴⁶ Yoshimi Yoshiaki, *Comfort Women: Sexual Slavery in the Japanese Military During World War II*, trans. Suzanne O'Brien (New York, 2000).

⁴⁷ Komagome Takeshi, *Shokuminchi teikoku Nihon no bunka tōgō* [Cultural Integration in the Japanese Colonial Empire] (Tokyo, 1996).

remains uncertain. What matters now is maintaining the conditions under which historical claims can still be made in ways that others can test and challenge.

Burnett's and Martinsen's experiments show that such a defence can be pedagogically viable. Students can learn to distinguish simulation from interpretation, but only through composition, only through experiencing friction that reveals where explanation fails. By requiring them to confront the failures of machine-generated prose, these assignments transform writing into a site of interrogative reflection. Students learn not only what to write but how to recognize when interpretation has occurred and when it has been simulated. They discover that understanding is not transmitted through fluent sentences but constructed through the resistance of writing, through the paragraph that will not settle, the transition that exposes contradiction, the revision that forces reconceptualization. This discovery cannot be automated. It requires the presence of a subject who judges, hesitates and revises under pressure.

If historical writing ceases to be legible as a record of interpretive labour, the discipline loses its capacity to produce knowledge that is both accountable and contestable. What will remain is prose that simulates historical thinking without enacting it, persuasive sequences that circulate without being inspectable. History as an academic practice will not vanish, but its authority and distinctiveness will be eroded as its methods become, in public perception, indistinguishable from stochastic simulations. Yet the very arrival of stochastic history constitutes a rupture the discipline can seize: By threatening to dissolve it, these systems have made visible the relationship between composition and knowledge that had previously required no defence.

Peer Review

For transparency, the peer review documents associated with this article are available at <https://doi.org/10.1111/1468-229x.70092>.