

Books

How AI is changing language



As allegations of LLM use rock the literary and media worlds, linguists explain what really distinguishes human and machine writing, while novelists including Jennifer Egan and Jeanette Winterson reflect on the future of fiction in an age of ChatGPT

Three paragraphs, from three different hotel reviews. Can you tell which, if any, were AI-generated?

“The hotel is in a great location for everything. Lots of places to eat and drink. The hotel itself is always abuzz. The tavern located on the ground floor is definitely a must. Food, service, prices and atmosphere were great.”

“A good hotel, though the room had the proportions of a well-appointed lift. Slept well, shower was excellent, staff were friendly. Breakfast was busy but competent. Would return, though probably not with a very large suitcase.”

“Excellent base for a London trip. The room was quiet, the bed comfortable, and everything worked exactly as it should. Staff were helpful without hovering. A smooth, unfussy stay from start to finish.”

How do you reckon you did? Most people, says Claire Hardaker, a professor of forensic linguistics at the University of Lancaster, get this kind of judgment right only about 60% of the time. Her online test, [Bot or Not](#), asks users to identify the fakes in a series of 15 reviews. The middling success rate might come as a surprise to those convinced they can spot AI writing at 50 paces. When [doubts](#) were raised in May about the authenticity of a prizewinning short story by Jamir Nazir, social media users were lightning-quick in their condemnation. “If you know, you know,” commented one. Nazir later [told](#) The Atlantic that he didn’t use AI.

Hardaker says her respondents tend to rely on a few quick rules of thumb to identify AI language, including the presence of clichés and the use of dashes. The “rule of three”, where words or phrases are arranged in a satisfying trio, is also thought to be a giveaway. “People have learned very simplistic rubrics and now just madly apply them everywhere.”

There’s a problem, though: these “tells” are also characteristic of human writing, which, after all, the large language models (LLMs) that produce them

were trained on. “You could go back to Charles Dickens and say he had AI, because he used the em dash too.” And orators have known about the rule of three ever since Julius Caesar said *Veni, vidi, vici*. In our hotel review examples, only the first one was authentic. Did you clock it?


Perhaps because it is so hard to know for sure, suspicion has become the order of the day. In the literary world, accusations of AI use now bedevil writers, with varying levels of justification. A debut horror novel, *Shy Girl*, was withdrawn by publishers Hachette after rumours circulated online that the author had relied on AI, which she denies; Steven Rosenbaum’s book *The Future of Truth*, a serious study of “how AI reshapes reality”, was found to contain numerous hallucinated quotations, which the author acknowledged in an apology.

Media organisations, including the Guardian, field increasing numbers of complaints about supposedly AI-generated text. These include intuitions about particular turns of phrase, but also comments about typos and grammatical errors. In one case, the word “after” was inadvertently duplicated in a sentence. “I can’t imagine a human editor/proofreader missing something like this,” wrote one reader, displaying a touching faith in our copy-editing abilities.


The problem is that not only does AI train on human writing, but humans are stylistically influenced by AI, the interplay creating a kind of linguistic hall of mirrors. Short of an author admitting it, it’s hard to say for certain whether an individual piece of writing is AI or not. That uncertainty is a recipe for paranoia.

And if you’re tempted to reach for a commercial screening tool to sort human from machine, that comes with uncertainty too, says Hardaker. “Given that some of us naturally write in a way that would be seen as AI-like” - she mentions neurodivergent people, for example - “that will be detected as AI. And you can modify AI output to make it seem more human-like. You put that kind of content into an AI detector, you’re going to get wacky results.” As someone who has served as an expert witness in court, she’s “extremely sceptical” about their efficacy.

The newly popular detector Pangram, which boasts false positive rates of around 1 in 10,000, has been shown in independent tests to be highly effective at detecting AI writing, even when it's been run through a "humanizer" app to disguise its origin. But questions remain. I was able to fool it on the first attempt (see the screenshot below) by channelling a bombastic register that might well be characteristic of AI, but could equally be the work of someone with a naturally bombastic style - or, more to the point, a writer who has been steeped in the output of the LLMs that power ChatGPT, Claude and Gemini. That, increasingly, is all of us.



token of human expression to destruction?



📷 Photograph: David Shariatmadari

Vast amounts of AI writing are now being published every day - from advertising copy to academic abstracts and fiction. At the same time, it looms ever larger over our lives via auto-generated email suggestions, “AI overview” search results, and the responses to our chatbot queries. At this level of exposure, it’s no longer a question of whether AI is changing language, both the way we speak and the way we write; the question is how. And should we resist, or embrace it?

We’ve known for some time that LLMs generate text that can be slightly different from human writing, *on average*. Often this only becomes clear when you look at large amounts of material. One eagle-eyed researcher linked the sudden popularity of the word “delve” to LLMs back in 2024 after searching a database of scientific papers. Other “focal words” that AIs have

tended to overuse include “showcase”, “boast”, “underscore”, “garner”, “align”, “surpass” and “intricate”. But, again, any individual piece of writing could entirely innocently make use of this vocabulary.

In a further twist, some researchers think the “delve” phenomenon might not be down to the models themselves, but the humans tasked with evaluating and steering them in a process known as “reinforcement learning with human feedback”. For workers who are “underpaid, stressed, and under time pressure”, it seems “certain words are treated as a proxy for quality” and the model is inadvertently trained to use them more often. In other words, “delve” might owe its meteoric rise to the fact it *doesn’t* seem like the kind of word an AI would use. (A separate suggestion that it appeared more often because it was characteristic of English used in Nigeria, where many RLHF workers lived, isn’t borne out by the data.)

There are other patterns we can distinguish: LLMs love nouns, but they seem to use pronouns less than humans. This might reflect the fact they don’t do as much talking about themselves or other people as we social creatures do. They like attributive adjectives (“the uncomfortable chair”), but not predicative ones (“the chair was uncomfortable”), perhaps because they prefer to deliver information in small, dense packages, whereas we pad things out. Different models have clear idiosyncrasies - you might even call them “dialects”: Gemini enjoys saying “here’s a breakdown”, while Deepseek often responds with a cheerful “Certainly!”. When asked to edit formal English from around the world, AI tends to flatten and homogenise towards an Anglo-American standard, in a process researchers have termed “cultural ghosting”. Thus the perfectly acceptable request in Indian professional English to “Kindly do the needful & revert back at the earliest” gets “corrected” to “Please complete the task & respond promptly.”

The evidence that aspects of LLM-speak have escaped into the “real” world, changing the way humans use language when AIs aren’t around, is now rolling in. One study analysed thousands of unscripted conversations and found that words like “delve” and “boast” spiked after ChatGPT was released. Another showed the frequency of “delve” in academic abstracts actually dropping after it was singled out on social media, in a sign that AI’s

influence might play out in complex ways.

Does any of this matter? Language changes all the time - words come in and out of fashion, and new technology has always been one of the forces behind this. AI does seem to be generating particularly high levels of anxiety, though. Why? “I think where it scares people is that idea of encroaching into sentience, of becoming the new human,” Hardaker says. Since 2023, she’s expanded the Bot or Not project into speech and music, and has noticed just how viscerally people react when a song they’ve enjoyed turns out to have been composed and performed by a machine.

Gary Shteyngart, a novelist who teaches creative writing at Columbia University, noticed a similar strength of feeling among his students at the prospect of AI literature. “When one of my graduate students said ‘as an experiment, I’m going to be writing a part of this piece with AI’, the other students became so angry, they wrote letters to me saying how awful this was.”

“There’s a kind of implicit bargain between writer and reader where you know the work that you’re getting is generated by a human being, and I think it felt like an assault on that,” he says. “Reading literary fiction is this incredible Vulcan mind meld with another human being, entering someone else’s consciousness. With AI I’m entering the simulacrum of another person’s consciousness, one degree removed, or many degrees removed. How sad is that by comparison?”

For Hardaker, “I guess it impinges on what we think of as what makes us special, what makes us valuable and unique”. At the same time, the music-generation model she uses “has generated some absolute bangers. I listen to them, unironically, in my car, and I enjoy them quite a lot.”

Could the same happen with literature? Will a machine-authored novel one day take its place among the 100 greatest of all time? Peter Stockwell, professor of literary linguistics at the University of Nottingham, thinks AI

may be able to do the basics, but it can't scale the heights. "If you want something that's very familiar and very mediocre and entirely functional, it's amazingly good at that."

One way to think of language, he says, is as a series of layers, with words at the bottom followed by phrases, clauses, compound sentences, all the way up to narrative structure. "AI is really good at the lower levels. It's learned lots of our syntactic structures and so everything looks well formed and grammatical. But, the higher up you go, the less good it is." The arc of a story is particularly hard for AI to get convincingly right.

"If you've got an AI to write a narrative, it can do a pretty good job of having a sequence of events and something happen at the end. But it wouldn't be a very *tellable* narrative," he continues. "Nothing startling or interesting would happen. And if there is anything startling, it will generally look like a mistake, rather than a brilliant twist."

The secret sauce of great writing remains secret - even to the academics who study it. "Linguists don't understand, really, how language works at its higher levels," at the level of discourse, storytelling, enchantment. "We can't build a machine to do something when we don't know how it works." We do have some idea of what it might boil down to - and that's our fundamentally social natures and, tied in with that, the fact that we are "wetware" - human flesh, with its spikes of adrenaline, rushes of dopamine, craving for social contact, all of which find expression in language's structure and the way we use it.

■ ■ *Would I work with an LLM? Of course! Why not?*

Jeannette Winterson

There are two broad models in linguistics, explains Stockwell, one that sees the brain as a computer, parsing grammatical structures and computing meaning from them, and the other that sees it as fundamentally embodied, something reflected in language by the fact that, in many languages, we understand by "seeing" or tend to think of "up", where our head is, as good (we get "high" and feel "low"). "One of the key things is that the current AIs don't have a body, they don't exist in the world, so they don't know what it

feels like to be in the world as a human.”

For Shteyngart, feeling is essential: “Today is the first warm day in New York. And if I was to start writing a novel, I think that [it] would be warmer. I think I would filter what I know through the warmth of the day. I think if I ate a really wonderful lunch and sat down to write, there would be more sensuousness in my writing.”

“The love of the body and its encounters with the physical world are what drives some of the best of literature. So I almost feel sorry for these LLMs, as I’m talking about them, because they’re pushed into some horrible machine in the Bay Area, and they just don’t know how wonderful life is.”

One much feared effect of the mass use of LLMs is that they act as a flattening force – smoothing away the variety and idiosyncrasy of human language into a kind of beige goo. That’s a legitimate concern, as far as it goes, though it’s not a new one. People have long angsted about the homogenising effects of American film and television on accent and vocabulary, and there are subgenres of language – political euphemism, customer-service prattle, therapy-speak – that have spread further from their home territory than many might like. The crucial thing, though, is that their influence tends to generate a backlash – and there’s no reason to think things will be different this time.

In fact, our capacity for innovation might ultimately be the thing that truly distinguishes human writing – particularly the literary kind – from AI. “The whole point of an LLM is that it’s trained on existing language. So it’s always retro,” says Stockwell. “I could get an AI and say ‘write me a short story in the style of Virginia Woolf’ and it’ll do a decent job. But what you can’t say is ‘write me a story in the unique style of the next great, serious literary innovator’. It couldn’t possibly do that.”

That’s because, once again, it lacks the social environment, and the body, that give rise to characteristically human motivations. “Why does somebody

do something new in an art form like literary writing? It can be out of annoyance or irritation with what's gone before. Or it's because somebody sees things in a different way than the run of the mill, or sometimes just because people are antsy and wanting to do something different, or a little bit crazy or isolated."

There are plenty of examples from history, says Stockwell: "After the bureaucracy and uniformity of the first world war, you've got this sudden, huge, antithetical artistic movement in the rise of surrealism and Dada; similarly, after the austerity of the second world war you get the psychedelic movement, art and literature changes again, quite radically. So there always seems to be that sort of kicking against the norms. It's hard to think how you would program an AI to do that, because AI works on an existing large body of material. It's the embodiment of the conservative-with-a-small-c status quo."

Originality is so important for novelist Jennifer Egan that she's quarantined herself from the technology entirely. "I feel a danger of infection, to use a kind of loaded metaphor," she tells me. "I know they stole some [of my] stuff, and there's nothing I can do about that, but I'm not giving them one more word voluntarily." Anthropic used pirated copies of books, including Egan's novels, to train its chatbot Claude; most LLMs use language from individual queries as additional training data. She sounds exasperated: "I don't want to partake of this kind of language spam that they're offering."

The zero-tolerance policy doesn't stop her from becoming paranoid. "I've been told a couple of stylistic things that are tells of AI, and they happen to be things I like. For example, I love em dashes, but I now find myself interrogating every one way more than I used to. I've also noticed that I'm prone to collections of three. So I find myself interrogating those as well. I don't mind that, actually, because the entire point is to not write something that anyone else could have."

What kind of advice would she give a younger writer now swimming in this water? Should they practice their own kind of AI hygiene? "I'm gonna now sound like the totally generic boomer that AI could probably have written, and my advice is: stay the fuck away. I mean, OK, use it to write emails. Even

use it to get research ideas. But if you want to be a writer: learn to write. Come on. I would really question why the impulse would be there to use it.”

Not everyone is so abstemious. Jeannette Winterson, who has written extensively about AI and art, tells me: “Every writer can make their own choice. Humans are tool-using animals. That has been our success story. At present all AI, including generative AI, is a tool. Would I work with an LLM? Of course! Why not?”

But she cautions against the view that AI’s linguistic competence means that it can equal or exceed human expression. “Beyond the basics, meaning becomes a series of inner realities and language is wonderful at conveying those inner realities. Machines do not share our reality, not least because they don’t have a limbic system. Humans cannot have a thought without a feeling ... literature is brilliant at revealing these layers.”

As I paste her quotes into a Google doc full of my notes, I notice the inbuilt AI making a suggestion. It asks whether I want to change Winterson’s words to more closely “match the style” of the existing material: to smooth over the edges of one of the English language’s most idiosyncratic writers. With an almost superstitious haste, I dismiss the prompt.

Most viewed
