

Chapter 13

AI in Communities at Risk and Housing Stability



Jong Cheol Shin and Kirk A. Foster

1 Introduction

Housing instability, broadly defined to include homelessness, overcrowding, substandard conditions, frequent residential moves, and significant cost burdens (Keene & Blankenship, 2023), remains a critical societal challenge, affecting millions globally and disproportionately impacting specific vulnerable populations. Various factors—such as poverty, insufficient affordable housing supply, gentrification, discriminatory redevelopment policies, and systemic inequities—combine to create precarious living conditions, ranging from struggles to afford rent to prolonged periods of homelessness. Within this multifaceted issue, “communities at risk” encompass diverse groups that are particularly vulnerable. These groups include, but are not limited to, low-income households facing affordability challenges, residents living in substandard or disaster-prone housing, racial and ethnic minorities confronting historical and ongoing discrimination, immigrants navigating unfamiliar or exclusionary housing systems, individuals with disabilities requiring tailored accommodations, veterans experiencing unique re-entry barriers, and populations currently experiencing homelessness. The intersection of economic hardship, environmental risks, and structural inequities disproportionately exacerbates housing precarity among these communities. Therefore, housing stability signifies more than mere shelter; it encompasses secure tenure, affordability, safe and healthy living conditions, and equitable access to essential community resources.

Jong Cheol Shin (✉)

Department of Population Health, Lehigh University, Bethlehem, PA, USA

e-mail: jos322@lehigh.edu

K. A. Foster

School of Social Work, The University of Texas at Arlington, Arlington, TX, USA

e-mail: kirk.foster@uta.edu

© The Author(s) 2026

R. An, M. A. Lindsey (eds.), *Artificial Intelligence in Social Work*,
https://doi.org/10.1007/978-3-032-18443-6_13

267

Housing instability is also deeply associated with a variety of public health outcomes. Chronic housing insecurity acts as a significant stressor, progressively impairing individuals' self-rated health the longer instability persists (Bhat et al., 2022). This chronic stress elevates the risk for severe health conditions—including hypertension, obesity, and diabetes—through mechanisms such as increased financial constraints limiting access to health-promoting resources, forced relocation leading to disadvantaged living environments, and heightened psychological stress and mental health challenges directly resulting from housing instability (Gu et al., 2023). Moreover, unstable housing negatively impacts chronic disease management, increasing disease prevalence among housing-insecure populations and complicating symptom control and treatment adherence (Nyamathi & Salem, 2021). Beyond direct health implications, housing instability is associated with increased vulnerability to violence, disrupted educational attainment and employment stability, and perpetuated intergenerational cycles of poverty (Olawade et al., 2023). Furthermore, poor-quality housing—a frequently overlooked aspect of housing insecurity—has direct links to respiratory illnesses, lead poisoning, injuries, and deteriorating mental health (Bashir, 2022). Exposure to hazardous environmental conditions due to substandard housing intensifies long-term health disparities (Anyanwu & Beyer, 2024). Thus, prolonged housing instability is not merely a socioeconomic hardship but also a critical determinant of public health, underscoring the importance of housing security and community well-being.

In recent years, artificial intelligence (AI)—encompassing methodologies such as machine learning, predictive analytics, natural language processing (NLP), and computer vision—has emerged as a promising set of tools to address aspects of the housing stability crisis. Advocates argue that AI can increase efficiency in resource management, yield data-driven insights into complex social patterns, and optimize the allocation of housing support and services. The appeal lies in the potential to leverage computational capacity to identify risks earlier, target interventions more accurately, and reduce the administrative burdens often experienced by service providers and recipients alike. For example, AI might be used to predict individual homelessness risk, expedite tenant screening, optimize the distribution of housing vouchers based on nuanced needs assessments, or monitor neighborhood conditions to guide targeted investments.

Nevertheless, deploying AI in the context of housing stability and vulnerable populations necessitates careful, critical examination. While AI offers potential operational advantages, its application carries profound ethical concerns. These include the risk not only of reflecting but also of amplifying existing societal inequities, creating new forms of exclusion, compromising individual privacy, and undermining critical principles such as fairness, transparency, and accountability. Algorithms, which often appear objective, can inadvertently mask biases and value judgments, making discriminatory or harmful outcomes harder to detect and challenge. Thus, this chapter advocates a critical, cautious, and fundamentally human-centered approach to AI, emphasizing equity and ethical considerations in its design, deployment, and governance within the housing sector.

The following sections provide an overview of primary AI techniques applied or proposed within housing contexts, detailing their functions, potential applications, and inherent limitations. Subsequently, the chapter explores specific applications of these technologies, assessing documented and anticipated impacts on various at-risk populations. Next, cross-cutting ethical, social, and governance challenges—such as algorithmic bias, transparency shortcomings, data privacy risks, surveillance concerns, the need for human oversight, and the importance of meaningful community participation in technology development—are examined in detail. Finally, the chapter concludes by summarizing key findings, acknowledging limitations, and offering recommendations for future research directions, policy development, and responsible AI practices. The overarching aim is to cultivate a nuanced understanding of AI's role in promoting housing stability—moving beyond simplified narratives of technological solutionism toward frameworks rooted in social justice, ethical accountability, and human dignity.

2 Overview of AI Techniques in Housing Stability

This chapter provides a foundational understanding of the key artificial intelligence methodologies relevant to housing stability challenges. The aim is an exhaustive technical exposition of these techniques and their potential applications, while simultaneously introducing the inherent limitations and risks that warrant critical scrutiny, particularly when deployed in contexts affecting vulnerable populations.

2.1 *Predictive Analytics and Machine Learning*

Predictive analytics involves using statistical algorithms and machine learning techniques to analyze historical and current data to forecast future events or outcomes. Machine learning (ML), a subset of AI, enables computer systems to learn patterns from data without being explicitly programmed for each task. Common techniques include regression analysis (predicting continuous values, such as rent prices or the duration of homelessness) and classification algorithms (assigning data points to predefined categories, such as “high risk” vs. “low risk” of eviction).

In the housing domain, these techniques are envisioned for applications such as early identification of individuals or families deemed at high risk of experiencing homelessness. By analyzing patterns in historical data, potentially including shelter usage history, interactions with social services, hospital visits, or criminal justice involvement, models aim to predict who might need preventive support. For example, a jurisdiction might use an ML model to flag at-risk households so that outreach teams can intervene with rental assistance or social support before an eviction or housing loss occurs (Byrne et al., 2019). Similarly, predictive analytics could

forecast neighborhood trends related to housing market changes; for instance, identifying areas potentially facing rapid gentrification or decline, ostensibly to inform urban planning or resource allocation decisions (Robb et al., 2022). One recent study demonstrated an AI model that could recognize early visual indicators of gentrification from Google Street View images, correctly predicting areas that would later undergo gentrification about 74% of the time (Huang et al., 2022). The potential benefit of such predictive approaches lies in moving from reactive crisis management to proactive intervention, thereby optimizing the use of limited preventive resources.

However, the efficacy and ethics of predictive systems are critically dependent on the data they consume. If the historical data reflects systemic biases, the models trained on it will inevitably learn and replicate those biases. For example, data drawn from criminal justice or emergency health services possibly over-represent marginalized communities due to factors like targeted policing or unequal access to primary healthcare, rather than reflecting any inherent predisposition to housing instability (Shaw et al., 2025). Consequently, predictive models might disproportionately flag individuals from these groups not because they are intrinsically more likely to become homeless, but because they are more visible to the systems generating the data. This act of forecasting future instability, based on a past shaped by systemic injustice, risks amplifying those injustices. In short, predictive modeling may improve efficiency, but without careful design and oversight, it can also entrench pre-existing biases under the guise of objectivity.

2.2 Algorithmic Scoring and Ranking

Building on predictive analytics, many housing-related systems employ algorithmic scoring and ranking to inform decisions. Tenant screening represents a widespread use of algorithmic scoring in the housing sector. Automated systems assess potential tenants based on factors like credit scores, income verification, eviction history, and criminal background checks, generating a score or recommendation (e.g., “accept,” “accept with conditions,” or “deny”) for each applicant, which landlords or property managers use to make leasing decisions. Similarly, algorithms are employed by housing authorities or social service agencies to prioritize applicants for scarce resources, such as housing vouchers, subsidized units, or supportive services. These systems often calculate a ‘vulnerability’ or ‘need’ score based on various indicators, ranking applicants to determine their place on waiting lists or eligibility for specific programs. The purported advantages are increased speed, consistency, and objectivity compared to manual review processes.

Despite the appeal of efficiency and objectivity, algorithmic scoring systems are fraught with potential pitfalls. By design, they reduce complex human circumstances and multifaceted needs to a single (often opaque) numerical value. The choice of which input variables to include and how to weigh them embeds specific

value judgments about what factors matter most—judgments that can inadvertently introduce or amplify bias. For instance, credit scores and eviction histories are often strongly correlated with race and income due to historical redlining, discriminatory lending practices, and economic inequality (Linde et al., 2023; Schneider, 2020). Criminal history and race reflect significant disparities in rental experiences, which are evident even among exonerees when they attempt to secure housing after release (Zannella et al., 2020). Using these factors heavily in a tenant screening score effectively encodes societal biases into the algorithm, meaning the resulting ‘objective’ score can produce discriminatory outcomes in practice.

A similar concern applies to scoring systems used for resource allocation: everything hinges on how concepts like “vulnerability” or “deservingness” are quantified. If an algorithmic formula prioritizes easily quantifiable indicators or those already captured in administrative databases, it may fail to account for more complex, intersectional needs. In one notable study on homelessness services, researchers developed a machine learning allocation system intended to minimize returns to homelessness (Kube et al., 2019); while it improved overall efficiency, they found that a significant fraction of households were algorithmically assigned to interventions that increased their risk of returning to homelessness, highlighting difficult trade-offs between efficiency and equity.

2.3 *Natural Language Processing (NLP)*

NLP is a branch of AI focused on enabling computers to understand, interpret, and generate human language. NLP techniques include methods like sentiment analysis (determining the emotional tone or attitude in text), topic modeling (discovering key themes in a large collection of documents), named entity recognition (extracting specific information such as names, locations, or dates from text), and machine translation (automatically translating text from one language to another), among others.

In the housing context, NLP can be applied to analyze large volumes of text that were previously difficult to review systematically. For example, property management companies or public agencies might use NLP tools to examine tenant communications—including emails, maintenance requests, complaint logs, and online reviews—to identify patterns and actionable insights. Furthermore, NLP methods have been used to develop models that identify overlooked case law regarding housing accessibility and eviction issues from extensive legal databases (Mohammadi et al., 2025). These applications facilitate a more comprehensive analysis of human rights jurisprudence related to housing. An NLP model might automatically flag recurring maintenance issues across multiple units or buildings, detect emerging disputes between tenants and landlords, gauge overall tenant sentiment or satisfaction, or even predict lease violations or the risk of eviction based on the content and tone of communications. Another potential application involves NLP-powered

chatbots to handle routine inquiries—answering frequently asked questions about housing applications, eligibility criteria, lease terms, renters’ rights, or available services. Such chatbots could provide quick, 24/7 responses and reduce the workload on housing office staff, theoretically improving responsiveness for tenants and applicants.

While NLP offers opportunities to extract insights from unstructured text data, its application in sensitive areas such as tenant communications raises significant concerns. The accuracy of NLP interpretation can be limited, particularly in nuance, sarcasm, language variations across cultures, and context-specific jargon, potentially leading to misinterpretations that could negatively affect tenants. More fundamentally, the systematic analysis of tenant communications constitutes surveillance, raising serious privacy concerns. Tenants may feel inhibited from raising legitimate concerns or complaints if they know their language is being algorithmically monitored and potentially used to assess them as “problematic.” This chilling effect undermines open communication, which is vital for healthy landlord-tenant relationships and timely issue resolution. The collection and analysis of such personal communication data require robust safeguards against misuse and breaches, particularly given the sensitive nature of housing-related issues. Therefore, any benefits in efficiency must be weighed against the erosion of privacy and trust between tenants and housing providers.

2.4 *Computer Vision*

Computer vision involves training AI systems to interpret and understand information from digital images or videos. Techniques include object detection (identifying specific objects within an image), image classification (categorizing images based on content), and scene analysis (interpreting the context and relationships within an image).

In urban studies and housing policy, computer vision has been explored as a tool for analyzing large-scale visual datasets, such as satellite imagery, aerial photos, and publicly available street-level images (e.g., Google Street View). Algorithms can be trained to identify visual indicators associated with neighborhood conditions—such as building upkeep (are roofs and exteriors intact or in disrepair?), the presence of litter or graffiti (often labeled as signs of “blight”), the amount of green space or tree cover, the condition of roads and sidewalks, the presence of abandoned cars, and so on. For instance, a study used pre- and post-disaster satellite images to assess hurricane damage in Florida, a task that might be difficult to achieve solely through manual surveys (Weber & Kané, 2020). Also, a city could use computer vision to scan current satellite imagery and detect properties with specific risk factors (e.g., tarp-covered roofs or structural damage) to prioritize housing code inspections or repairs (Deng et al., 2022). Researchers used a combination of Google Street View imagery and satellite data, along with CNNs, to estimate property

values in London. Features such as garden size and building condition explained 68% of the price variance, aiding equitable tax assessments (Law et al., 2019). This kind of urban visual intelligence could provide planners and policymakers with broad, data-driven insights into the physical environment relevant to housing stability and community well-being.

Despite these promising applications, relying on computer vision for neighborhood assessment carries inherent risks and limitations. Visual data alone provides an incomplete picture of community well-being, often overlooking crucial social factors such as community cohesion, informal support networks, and access to services—factors not visible in an image. Moreover, the interpretation of visual cues like “blight” or “disrepair” can be highly subjective and culturally biased. Algorithms trained on datasets labeled according to dominant aesthetic norms might unfairly penalize neighborhoods with different cultural expressions or predominantly lower-income areas, whose properties may look different but still provide stable homes. Such algorithmic assessments risk automating and scaling subjective judgments, potentially reinforcing spatial stigma and justifying disinvestment or gentrification pressures in areas algorithmically labeled as ‘declining’ based on superficial visual indicators. Without careful validation, ground-truthing, and consideration of context, computer vision analysis of neighborhoods could perpetuate, rather than challenge, spatial inequalities.

2.5 Reinforcement Learning (RL)

RL is a machine learning approach in which an AI agent learns optimal decision-making through environmental interaction and feedback, including rewards and penalties. Unlike traditional models using static datasets, RL employs a dynamic trial-and-error methodology: the agent acts, evaluates the outcome (including a reward signal indicating relative success), and progressively refines its strategy to maximize cumulative rewards. This approach has demonstrated efficacy across complex sequential decision-making tasks, including gaming, robotics, and resource allocation.

In the context of housing stability, while RL applications remain nascent, the potential is significant. Housing systems feature dynamic interactions—individuals navigating shelter systems, program applications, and market responses to policy changes—making them theoretically suitable for RL optimization. An RL model could be trained within agent-based simulations of eviction-prevention programs or homelessness systems to identify optimal intervention policies that minimize housing instability or improve resource allocation efficiency (Sarkar et al., 2024). Recent research in the UK incorporated RL algorithms into housing market simulations to analyze how the market adapts to economic disruptions (Olmez et al., 2024). Additionally, RL has demonstrated the capacity to design more equitable resource distribution frameworks, as evidenced by its application to optimize the geographic

placement of refugees to enhance employment outcomes (Bansak et al., 2018). RL's ability to balance immediate and long-term rewards aligns well with policy planning requirements.

Nevertheless, RL implementation presents substantial challenges. The reward function, which defines success, requires meticulous design, as imprecise specifications can lead to unintended consequences. Furthermore, effective RL training necessitates simulation environments that accurately reflect real-world conditions, which proves especially difficult for complex social systems like housing. Biased or incomplete simulation data will inevitably lead to flawed RL policies. Finally, the decision-making processes of RL systems, similar to other sophisticated AI models, often lack transparency, potentially complicating accountability when determining why specific allocation strategies are recommended.

3 AI Applications and Differential Impacts by Population

Moving from general AI capabilities to their concrete use in the housing sector, this section examines how these technologies are being applied (or proposed) in real-world scenarios, with a focus on impacts across various at-risk populations and communities.

3.1 Homelessness Prevention and Support

One of the most prominent emerging applications of AI in this domain is the use of predictive models to prevent homelessness. AI-driven predictive analytics systems are being developed and deployed to identify individuals and families at high risk of homelessness, enabling preventive interventions before a crisis hits. These systems typically analyze large administrative datasets, drawing on a combination of data sources, including homeless shelter intakes, subsidized housing waitlists, interactions with social services, healthcare utilization (e.g., emergency room visits, mental health services), and, sometimes, encounters with the criminal justice system.

By detecting patterns (for example, a history of eviction filings or frequent address changes coupled with hospital visits), the model generates a risk score or flag for those it deems most likely to experience homelessness in the near future (Byrne et al., 2019; VanBerlo et al., 2021). Social service providers can then target outreach and support to high-risk individuals, offering case management, short-term financial assistance (such as help with rent or utility arrears), or referrals to programs, such as landlord mediation, before those individuals lose their current housing.

On the back end, such algorithms may also be used to optimize resource allocation for those already experiencing homelessness. For instance, if shelter beds or

transitional housing units are limited, an algorithm might prioritize who should receive these resources by predicting which individuals would benefit most or have nowhere else to go (Kube et al., 2023). These efforts can shift the homelessness response from reactive to proactive and use data to deploy scarce resources in the most effective way possible. In Los Angeles County, for example, a pilot AI program called Housing for Health has been used to predict which residents are most likely to lose their housing and then provide financial assistance and services to those individuals. Early reports indicate that since 2021, the program has enrolled over 700 people identified by the algorithm, and about 86% of them have retained their housing, suggesting some success in preventing new homelessness cases (Kendall, 2024). Beyond the resource relocation, the UCLA California Policy Lab used an AI-powered predictive analytics system to identify individuals at the highest risk of homelessness. This system achieved nearly 3.5 times the accuracy of trained social workers.

3.2 Natural Disaster Victims: AI for Disaster Risk Management and Resilient Housing

Communities that are prone to natural disasters (such as hurricanes, floods, wildfires, or earthquakes) face unique risks to housing stability. In recent years, AI has been increasingly applied to disaster risk management, with implications for protecting housing and communities in hazard-prone areas (Sun et al., 2020). These applications are especially relevant for residents of disaster-prone regions, who are often among the most vulnerable to sudden homelessness or long-term displacement when disaster strikes.

One major use of AI here is in predictive forecasting and early warning systems. Machine learning models can be trained on vast datasets of weather patterns, geological data, and historical disaster impacts to improve the prediction of events such as floods, storms, and landslides (Zhu et al., 2019). For example, AI models have been developed that analyze meteorological data in real time to predict the likelihood and severity of extreme weather events (such as heavy rainfall leading to floods, drought, etc.) more accurately and earlier than traditional models (Adikari et al., 2021). By providing more timely and precise warnings, these AI systems enable at-risk residents to evacuate or reinforce their homes sooner, potentially saving lives and reducing housing damage. Early warning can be critical for low-income or elderly residents who need more time to move to safety.

Another application is risk mapping: AI can integrate diverse data (topography, soil, drainage, building characteristics, and socioeconomic data) to identify which neighborhoods or buildings are most vulnerable to future disasters. This information helps officials target pre-disaster mitigation efforts, such as strengthening infrastructure or prioritizing buildings for retrofits. Researchers at Texas A&M, for example, have used AI and Google Street View imagery to assess flood risk at a very

local level, aiming to pinpoint community vulnerabilities that current broad flood maps might miss (Ho et al., 2024). By identifying that, the map also serves as a valuable supplement and reference to support forecasting extreme disasters, to shore up the community, or at least to ensure its residents have evacuation and recovery plans (Emami & Marzban, 2023).

AI is also being harnessed during and after disasters to aid response and recovery. Computer vision algorithms have been used on satellite and aerial imagery to rapidly assess damage after a disaster, for example, detecting which buildings are destroyed or uninhabitable in the aftermath of a flood or earthquake (Jing et al., 2022; Kim et al., 2022). This can significantly speed up damage assessments, allowing emergency managers to allocate shelter and repair resources to the hardest-hit areas without waiting for manual inspections of every structure. AI analysis of social media and emergency calls can likewise help first responders identify where people are trapped or in urgent need of help during an unfolding disaster (Imran et al., 2020). All these efforts contribute to stabilizing housing for survivors by getting aid where it is needed faster and moving affected individuals into stable housing.

3.3 Housing Service Applicant and Refugees: Resource Allocation for Housing Services

Housing and social service agencies often face the challenging task of allocating insufficient resources; housing vouchers, public housing units, and supportive program placements are consistently inadequate to meet the needs of all qualified applicants. Traditionally, allocation decisions—determining priority access and waitlist positioning—have relied on first-come-first-served approaches, categorical preferences (such as prioritizing families with children or veterans), and individual caseworker judgment. However, the sector is increasingly exploring AI and algorithmic solutions to manage resource distribution more systematically and to prioritize applicants based on assessed need or anticipated benefit (Chui et al., 2018; Yigitcanlar et al., 2020). Consider a public housing authority with 100 rental vouchers available for 1000 eligible applicants. An algorithmic ranking system could assign each applicant a comprehensive vulnerability score to determine assistance priority. This score might incorporate multiple factors, including current housing status (homeless, doubled-up, cost-burdened), income-to-poverty ratio, number of dependents, health conditions, disability status, and application longevity. In theory, this approach promotes consistency through standardized evaluation criteria and enhances efficiency by directing limited resources to the most urgent cases, potentially preventing further hardship.

Real-world implementations include coordinated entry systems for homelessness services that use standardized assessment tools to generate vulnerability scores that guide program placements (Kube et al., 2023). Some jurisdictions are

enhancing these systems with machine learning models that predict the specific interventions most likely to prevent individuals from returning to homelessness, allocating resources accordingly. During the COVID-19 pandemic, government entities such as the Colorado Division of Housing, faced with unprecedented demand for emergency rental assistance, employed algorithms to evaluate and prioritize applicants in cohorts, aiming to identify those with the greatest need and highest likelihood of successful outcomes while ensuring equitable distribution across neighborhoods and demographic groups (Chuang, 2021). AI applications also extend to refugee housing allocation, supporting the development of adequate low-income housing solutions to prevent slum formation (Marji & Kohout, 2023) while simultaneously enhancing employment opportunities (Bansak et al., 2018; Ahani et al., 2021).

3.4 Tenant Screening and Housing Access: AI Systems, Eviction Risk, and Demographic Disparities

The adoption of automated tenant screening software represents one of the most pervasive and consequential applications of artificial intelligence within the housing market. Property managers and landlords increasingly use these technological tools to evaluate prospective tenants more efficiently. These screening systems aggregate and analyze comprehensive data from multiple sources—credit reports, eviction histories in public records databases, criminal background checks, income verification services, and occasionally social media platforms or third-party data brokers. Upon processing this information, the software generates an applicant assessment in the form of a numerical score, a rating, or a categorical recommendation (e.g., approve, deny, or approve with additional conditions).

From the property owner's perspective, these automated systems offer significant operational advantages through accelerated processing capabilities and perceived objectivity—applications can be evaluated within minutes, ostensibly reducing the risk of leasing to individuals who might default on rental payments or create management challenges. However, the widespread implementation of algorithmic tenant screening potentially erects substantial barriers to housing accessibility, particularly affecting historically disadvantaged populations (Schneider, 2020).

Beyond individual tenant screening, artificial intelligence applications now extend to property-level eviction risk prediction by integrating diverse datasets, including spatiotemporal public records and municipal data. By synthesizing neighborhood characteristics, property ownership data, and historical eviction patterns, these predictive models enable social service professionals to target at-risk households more strategically, thereby enhancing intervention efficiency and effectiveness (Robb et al., 2022; Mashiat et al., 2024). Nevertheless, these eviction records and predictions can simultaneously function as mechanisms for applicant exclusion or tenant stigmatization. Research employing natural language processing to extract

and analyze nationwide eviction filings from court records has uncovered systematic inequities, including unlawful eviction proceedings and disproportionately elevated eviction rates among non-English-speaking tenant populations (Thomas et al., 2024).

4 Cross-Cutting Ethical, Social, and Governance Issues

The specific applications we have examined point toward several recurring, cross-cutting themes that demand deeper reflection. These themes relate to the fundamental nature of algorithmic bias, the challenges of transparency and accountability, data privacy and surveillance implications, the complexities of human-AI interaction in decision-making, and the critical role of power dynamics and community participation in shaping technological outcomes. In this section, we synthesize broader issues that cut across the specific use cases to understand the societal and ethical stakes involved in deploying AI in the sensitive domain of housing stability.

4.1 *Algorithmic Bias and Fairness Challenges*

A central theme emerging from the previous section is that algorithmic bias is not merely an occasional error but a systemic issue often deeply embedded in AI systems used in the housing context. This bias arises when AI systems learn from data that reflects existing societal inequities. From homelessness prediction to tenant screening, we have seen how algorithms trained on historical data that reflects society's inequities will inevitably reproduce those inequities. The data itself is not neutral, and it captures past and present discrimination in policing, lending, healthcare access, and housing patterns. Consequently, even if AI models do not explicitly use protected characteristics such as race, they can still produce discriminatory outcomes. This occurs through reliance on proxy variables—like zip codes, income levels, criminal records, or even names—that correlate strongly with sensitive traits. For example, a tenant screening algorithm might exclude race as a direct input. Yet, if it heavily weighs credit scores, criminal histories, and prior evictions—factors already influenced by systemic racial disparities—the algorithm will likely yield disproportionately harmful outcomes for minority or low-income applicants, effectively engaging in proxy discrimination.

Addressing such pervasive bias is a complex, twofold challenge: first, to minimize bias within AI systems themselves, and second, to address the societal biases AI observes. For the former, AI practitioners are developing technical “debiasing” methods, conducting bias audits before and during deployment (which involve examining error rates and outcomes across demographic groups), incorporating domain experts and affected communities in defining fairness criteria, and

implementing safeguards like removing specific inputs or using post-processing techniques to balance outcomes. Regulatory bodies are also increasingly considering algorithmic accountability laws that would mandate impact assessments for bias in high-stakes AI systems.

The latter challenge—eliminating bias from the world AI observes—is a broader societal endeavor. For communities at risk, it is crucial that any AI deployed in housing actively counteracts historical injustices rather than inadvertently perpetuating them. Achieving this requires robust interdisciplinary collaboration: technologists must work alongside substantive experts, legal experts, and the affected communities to define fair outcomes and continuously adapt systems toward these goals.

4.2 Transparency, Explainability, and Accountability

Another cross-cutting concern is the lack of transparency and explainability surrounding many AI systems used in housing decisions. Whether due to the inherent complexity of machine learning models (“black box” algorithms issue) or the protection of proprietary algorithms, the reasoning behind automated decisions is often opaque to those affected by them. This problem, in context, is that there is no explanation for tenant screening rejections or for resource allocation priorities. This opacity poses fundamental challenges to accountability and procedural justice in housing. When individuals cannot understand why a decision impacting their housing stability was made, they are effectively denied the ability to identify potential errors, contest unfair biases, or meaningfully exercise their right to appeal. This lack of transparency is not merely a technical inconvenience; it undermines fundamental principles of due process, particularly when basic needs are at stake.

Meaningful transparency requires not only open-sourcing code but also explanations understandable to non-experts that detail how specific inputs led to a particular outcome. Different levels of explanation may be needed for various audiences, ranging from understanding the system’s general logic to obtaining a specific rationale for an individual decision. Achieving this is not trivial, of course, especially for complex machine learning models that even their creators might struggle to interpret precisely.

Furthermore, even if explanations are provided, effective recourse requires accessible and timely mechanisms for appeal and correction. Without robust transparency and avenues for redress, AI systems risk operating as unaccountable arbiters of housing opportunity, disproportionately harming vulnerable individuals who often lack the resources or technical literacy to challenge opaque algorithmic decisions. This is particularly problematic for vulnerable individuals who often lack resources or technical literacy to navigate complicated systems. Suppose the computer’s logic is secret or unintelligible. In that case, it may be hard to hold anyone accountable and accept the denial as final, even if it was in error or unfair. The right

to shelter is so fundamental that any decision-making process must be seen as legitimate and fair.

4.3 Data Privacy, Surveillance, and Data Governance

The deployment of AI systems in the housing sector relies heavily on the collection and analysis of vast amounts of personal data. Whether this involves individual financial, health, and behavioral data for predictive modeling and tenant screening, or aggregated data on communities and properties for neighborhood analysis, the drive to enhance algorithmic performance often leads to more expansive and potentially invasive data-gathering practices. This trend raises significant concerns regarding privacy, informed consent, and effective data governance, particularly when applied to vulnerable populations who may have limited capacity to refuse data collection due to their need for assistance or housing.

Consider the types of data we discussed: homeless service databases may contain detailed personal histories, including interactions with shelters and outreach services, as well as information on mental health or substance use. Tenant screening processes access credit reports, criminal records, and eviction histories. Resource allocation algorithms might utilize health records or disability status. Furthermore, neighborhood surveillance for monitoring purposes can capture images or sensor data from public spaces that inadvertently record individuals' daily activities.

The aggregation of such extensive personal data into centralized systems creates substantial risks, including data breaches and misuse. There is also the prevalent risk of "function creep," where data collected for a specific, stated purpose is later repurposed for unrelated uses without individuals' explicit consent. Marginalized communities have historically been disproportionately subjected to such practices and heightened surveillance. These privacy risks are particularly acute for at-risk individuals because they often face a coercive dynamic. For example, accessing essential services like rental assistance may be contingent on signing broad data-release forms, effectively trading privacy for basic needs.

To mitigate these dangers, robust data governance frameworks are indispensable. Such frameworks should include clear rules on data minimization (collecting only what is strictly necessary for the defined purpose), purpose limitation (using data solely for its original purpose unless new, informed consent is obtained), and secure data storage and handling procedures to protect against breaches. Strict limitations on data sharing, especially between different agencies and with law enforcement (unless mandated by due legal process), are crucial to prevent function creep. Furthermore, individuals whose data is used in algorithmic decision-making should have the right to access, review, and correct their data, a principle closely tied to transparency and accountability. Without these comprehensive safeguards, the increasing datafication driven by AI adoption

could significantly erode privacy, expand surveillance, and further marginalize already vulnerable populations.

4.4 Human Agency, Oversight, and Automation Bias

Given the complexities and potential biases in AI decision-making within the housing sector, a frequently proposed safeguard is the integration of “human-in-the-loop” systems. This approach ensures that human judgment and oversight remain integral to AI-driven processes. Theoretically, human case workers, housing officers, or managers can complement algorithmic recommendations by applying contextual understanding, empathy, and the ability to identify errors or extenuating circumstances that an algorithm might overlook. For instance, if an algorithm deems an applicant ineligible for a housing program, a human reviewer could investigate further, potentially uncovering data errors or specific situations not accounted for by the model.

However, achieving meaningful human oversight in practice is challenging and not automatically guaranteed by simply including a human in the process. A significant hurdle is “automation bias”—the documented human tendency to over-trust or unduly defer to automated systems’ suggestions, even when these suggestions may be flawed. In high-pressure environments like busy housing offices, staff might treat algorithmic outputs (e.g., risk scores or eligibility recommendations) as definitive and approve them without thorough scrutiny, particularly if they face time constraints or staffing levels have been reduced under the assumption that AI is efficient. This can lead to an “automation paradox”: while a human is theoretically in the loop, their practical reliance on automation can diminish the value and effectiveness of their oversight.

For human-in-the-loop oversight to be genuinely effective, several conditions are essential. First, human reviewers require adequate training to understand the AI system’s functionality, strengths, and limitations, and the true implications of its outputs. Access to information is also critical; reviewers should understand the rationale behind an algorithm’s conclusion (linking back to transparency and explainability) and be able to consider supplementary details not included in the algorithm’s dataset. Furthermore, sufficient time and resources must be allocated for thoughtful review; perfunctory oversight due to excessive workloads will not be meaningful.

Crucially, human reviewers must possess the authority and organizational support to override algorithmic recommendations when appropriate, without fear of negative repercussions. If an organization discourages or penalizes staff for deviating from AI-generated suggestions, the oversight function becomes performative. Conversely, if overriding the algorithm is accepted and even encouraged in suitable situations, human judgment can serve as a vital check. Policies could mandate, for example, that any application denied by an algorithm undergoes automatic human

review before final rejection, or that humans randomly audit a certain percentage of all algorithmic decisions (approvals and denials) to detect errors.

In summary, while maintaining a human presence in AI-driven housing decisions is necessary, it is not, by itself, sufficient to ensure fairness and accuracy. The implementation must thoughtfully empower human judgment rather than marginalizing it. This requires careful design of the human-AI interface, comprehensive training programs, and clear protocols for decision-making and appeals to ensure that human agency is preserved and enhanced, not undermined by automation.

4.5 Power Asymmetries and Community Participation

A fundamental cross-cutting issue in the deployment of AI systems within the housing sector is the inherent power imbalance among stakeholders. Those typically driving AI adoption—technology companies, government agencies seeking efficiencies, or large property management firms—often possess significantly more power, resources, and technical expertise than the tenants, applicants, and community members whose lives are directly impacted by these systems. This asymmetry can shape which priorities are encoded into AI systems, whose values are reflected in their design, and which concerns are addressed and which are overlooked. Consequently, AI systems may be developed in a top-down manner, prioritizing institutional goals such as efficiency, cost savings, and risk management, potentially without adequate regard for residents' needs and rights, including secure housing, dignity, and autonomy.

Inclusive design principles and meaningful community participation are increasingly recognized as essential correctives to these power imbalances. The core idea is that individuals most affected by an AI system—for example, public housing residents subject to new monitoring technologies, or homeless individuals whose data informs risk prediction models—should have a significant, empowered voice throughout the technology's lifecycle. This includes problem identification, system design and testing, deployment, and ongoing governance and oversight. Such participation must extend beyond tokenistic gestures, such as superficial focus groups; it requires genuine power-sharing in decision-making.

Achieving meaningful community participation involves overcoming several challenges. Firstly, bridging the knowledge gap is crucial; communities may require resources and support to understand complex AI technologies and their potential impacts, enabling them to engage as informed and equal partners. This could involve public education initiatives or funding for community organizations to access independent technical expertise. Secondly, participation must be early and continuous, not merely an afterthought or consultation once a system is already developed; belated input often cannot substantively alter core design choices.

Moreover, community input must have tangible influence, leading to fundamental changes in system design or deployment decisions, rather than being a procedural checkbox. This may require agencies to be willing to halt or significantly modify AI deployments in response to strong community opposition or concerns. It

could also involve granting community oversight bodies the authority to audit AI systems or even mandate their discontinuation if they prove harmful.

This participatory approach is not only ethically imperative but can also lead to more effective and equitable outcomes by diversifying the perspectives involved in problem-solving and system design. Building trust is a key benefit; if communities perceive that AI systems are developed with their input and genuinely serve their needs, rather than primarily functioning as tools for surveillance or control, then acceptance and collaboration are more likely.

5 Conclusion

This chapter has critically examined AI's role in housing for at-risk communities, revealing a landscape where promised efficiencies often collide with significant ethical risks. Key findings highlight that AI systems, by ingesting data reflecting societal biases, can amplify inequities related to race, income, and disability. This can create new barriers to housing, unfairly target vulnerable groups, and undermine due process due to opaque algorithms. The extensive data required also raises serious privacy and surveillance concerns, particularly for marginalized populations.

The central argument is that a critical, human-centered, and equity-focused approach is vital. Technological solutions must be rigorously weighed against potential harm to fairness, justice, and human dignity, and inherently balanced with human knowledge and experience. This necessitates independent research on AI's real-world impacts, including bias audits and the development of context-aware fairness frameworks. Stronger regulatory oversight is needed, mandating transparency, accountability, non-discrimination, and robust data protection. Housing agencies and practitioners must prioritize ethical guidelines, ongoing monitoring for bias, staff training on AI limitations, and meaningful community participation in the deployment of technology.

Ultimately, AI is a tool; its value in addressing housing instability hinges on its alignment with principles of equity and justice. Without such commitment, AI risks entrenching disadvantages rather than fostering the right to safe, stable housing for all.

Competing Interests The authors have no conflicts of interest to declare that are relevant to the content of this chapter.

References

- Adikari, K. E., Shrestha, S., Ratnayake, D. T., Budhathoki, A., Mohanasundaram, S., & Dailey, M. N. (2021). Evaluation of artificial intelligence models for flood and drought forecasting in arid and tropical regions. *Environmental Modelling & Software*, 144, 105136.
- Ahani, N., Andersson, T., Martinello, A., Teytelboym, A., & Trapp, A. C. (2021). Placement optimization in refugee resettlement. *Operations Research*, 69(5), 1468–1486.

- Anyanwu, C., & Beyer, K. M. (2024). Intersections among housing, environmental conditions, and health equity: a conceptual model for environmental justice policy. *Social Sciences & Humanities Open*, 9, 100845.
- Bansak, K., Ferwerda, J., Hainmueller, J., Dillon, A., Hangartner, D., Lawrence, D., & Weinstein, J. (2018). Improving refugee integration through data-driven algorithmic assignment. *Science*, 359(6373), 325–329.
- Bashir, S. A. (2022). Home is where the harm is: Inadequate housing as a public health crisis: American Journal of Public Health, 2002. In *The Affordable Housing Reader* (pp. 123–129). Routledge.
- Bhat, A. C., Almeida, D. M., Fenelon, A., & Santos-Lozada, A. R. (2022). A longitudinal analysis of the relationship between housing insecurity and physical health among midlife and aging adults in the United States. *SSM-Population Health*, 18, 101128.
- Byrne, T., Montgomery, A. E., & Fargo, J. D. (2019). Predictive modeling of housing instability and homelessness in the Veterans Health Administration. *Health services research*, 54(1), 75–85.
- Chui, M., Harryson, M., Valley, S., Manyika, J., & Roberts, R. (2018). *Notes from the AI frontier applying AI for social good*.
- Chuang T. (2021, October 25) *Colorado's COVID rent-assistance program may have paid up to \$4.4 million in error; audit finds*. The Colorado Sun <https://coloradosun.com/2021/10/25/pandemic-rent-assistance-covid-error-auditor-landlord/>
- Deng, J., Singh, A., Zhou, Y., Lu, Y., & Lee, V. C. S. (2022). Review on computer vision-based crack detection and quantification methodologies for civil structures. *Construction and Building Materials*, 356, 129238.
- Emami, P., & Marzban, A. (2023). The synergy of artificial intelligence (AI) and geographic information systems (GIS) for enhanced disaster management: opportunities and challenges. *Disaster Medicine and Public Health Preparedness*, 17, e507.
- Ho, Y. H., Lee, C. C., Diaz, N., Brody, S., & Mostafavi, A. (2024). Elev-vision: Automated lowest floor elevation estimation from segmenting street view images. *ACM Journal on Computing and Sustainable Societies*, 2(2), 1–18.
- Gu, K. D., Faulkner, K. C., & Thorndike, A. N. (2023). Housing instability and cardiometabolic health in the United States: a narrative review of the literature. *BMC Public Health*, 23(1), 931.
- Huang, T., Dai, T., Wang, Z., Yoon, H., Sheng, H., Ng, A. Y., et al. (2022). Detecting neighborhood gentrification at scale via street-level visual data. In *2022 IEEE International Conference on Big Data (Big Data)* (pp. 1632–1640). IEEE.
- Imran, M., Ofli, F., Caragea, D., & Torralba, A. (2020). Using AI and social media multimodal content for disaster response and management: Opportunities, challenges, and future directions. *Information Processing & Management*, 57(5), 102261.
- Jing, Y., Ren, Y., Liu, Y., Wang, D., & Yu, L. (2022). Automatic extraction of damaged houses by earthquake based on improved YOLOv5: A case study in Yangbi. *Remote Sensing*, 14(2), 382.
- Keene, D. E., & Blankenship, K. M. (2023). The affordable rental housing crisis and population health equity: A multidimensional and multilevel framework. *Journal of Urban Health*, 100(6), 1212–1223.
- Kendall, M. (2024, March 6) *This California county is testing AI's ability to prevent homelessness* <https://calmatters.org/housing/homelessness/2024/03/california-homeless-los-angeles-ai/>
- Kim, D., Won, J., Lee, E., Park, K. R., Kim, J., Park, S., et al. (2022). Disaster assessment using computer vision and satellite imagery: Applications in detecting water-related building damages. *Frontiers in Environmental Science*, 10, 969758.
- Kube, A., Das, S., & Fowler, P. J. (2019). Allocating interventions based on predicted outcomes: A case study on homelessness services. In *Proceedings of the AAAI Conference on Artificial Intelligence* (Vol. 33, No. 01, pp. 622–629).
- Kube, A. R., Das, S., & Fowler, P. J. (2023). Community-and data-driven homelessness prevention and service delivery: optimizing for equity. *Journal of the American Medical Informatics Association*, 30(6), 1032–1041.

- Law, S., Paige, B., & Russell, C. (2019). Take a look around: using street view and satellite images to estimate house prices. *ACM Transactions on Intelligent Systems and Technology (TIST)*, 10(5), 1–19.
- Linde, S., Walker, R. J., Campbell, J. A., & Egede, L. E. (2023). Historic residential redlining and present-day social determinants of health, home evictions, and food insecurity within US neighborhoods. *Journal of General Internal Medicine*, 38(15), 3321–3328.
- Marji, N., & Kohout, M. (2023). Optimizing the Spatial Organization of Refugee Camps in Jordan Through Artificial Intelligence. In *World Congress of Architects* (pp. 393–404). Springer International Publishing.
- Mashiat, T., DiChristofano, A., Fowler, P. J., & Das, S. (2024, June). Beyond eviction prediction: Leveraging local spatiotemporal public records to inform action. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency* (pp. 1383–1394).
- Mohammadi, M., Wieling, M., & Vols, M. (2025). An interpretable approach to detect case law on housing and eviction issues within the HUDOC database. *Artificial Intelligence and Law*, 1–22.
- Nyamathi, A. M., & Salem, B. E. (2021). The impact of unstable housing on health. *The Lancet Public Health*, 6(5), e265–e266.
- Olawade, D. B., Wada, O. J., David-Olawade, A. C., Kunonga, E., Abaire, O., & Ling, J. (2023). Using artificial intelligence to improve public health: a narrative review. *Frontiers in Public Health*, 11, 1196397.
- Olmez, S., Heppenstall, A., Ge, J., Elsenbroich, C., & Birks, D. (2024). Mitigating housing market shocks: an agentbased reinforcement learning approach with implications for real-time decision support. *Journal of Simulation*, 18(6), 921–939.
- Robb, K., Amigo, N. D., Marcoux, A., McAteer, M., & De Jong, J. (2022). Using integrated city data and machine learning to identify and intervene early on housing-related public health problems. *Journal of Public Health Management and Practice*, 28(2), E497–E505.
- Sarkar, A., DiChristofano, A., Das, S., Fowler, P. J., Jacobs, N., & Vorobeychik, Y. (2024, May). Geospatial active search for preventing evictions. In *Proceedings of the 23rd International Conference on Autonomous Agents and Multiagent Systems* (pp. 2456–2458).
- Schneider, V. (2020). Locked out by big data: how big data algorithms and machine learning may undermine housing justice. *Columbia Human Rights Law Review*, 52, 251.
- Shaw, C., Janeway, H., Preston-Suni, K., & Ryus, C. R. (2025). Exploring the complexity of homelessness in emergency medicine: Dissecting myths, evidence, and solutions. *AEM Education and Training*, 9, S108–S115.
- Sun, W., Bocchini, P., & Davison, B. D. (2020). Applications of artificial intelligence for disaster management. *Natural Hazards*, 103(3), 2631–2689.
- Thomas, T., Ramiller, A., Ren, C., & Toomet, O. (2024). Toward a national eviction data collection strategy using natural language processing. *Cityscape*, 26(1), 241–260.
- VanBerlo, B., Ross, M. A., Rivard, J., & Booker, R. (2021). Interpretable machine learning approaches to prediction of chronic homelessness. *Engineering Applications of Artificial Intelligence*, 102, 104243.
- Weber, E., & Kané, H. (2020). Building disaster damage assessment in satellite imagery with multi-temporal fusion. *arXiv*. <https://arxiv.org/abs/2004.05525>
- Yigitcanlar, T., Desouza, K. C., Butler, L., & Roozkhosh, F. (2020). Contributions and risks of artificial intelligence (AI) in building smarter cities: Insights from a systematic review of the literature. *Energies*, 13(6), 1473.
- Zannella, L., Clow, K., Rempel, E., Hamovitch, L., & Hall, V. (2020). The effects of race and criminal history on landlords' (un) willingness to rent to exonerees. *Law and Human Behavior*, 44(4), 300.
- Zhu, X., Zhang, G., & Sun, B. (2019). A comprehensive literature review of the demand forecasting methods of emergency resources from the perspective of artificial intelligence. *Natural Hazards*, 97, 65–82.

Open Access This chapter is licensed under the terms of the Creative Commons Attribution 4.0 International License (<http://creativecommons.org/licenses/by/4.0/>), which permits use, sharing, adaptation, distribution and reproduction in any medium or format, as long as you give appropriate credit to the original author(s) and the source, provide a link to the Creative Commons license and indicate if changes were made.

The images or other third party material in this chapter are included in the chapter's Creative Commons license, unless indicated otherwise in a credit line to the material. If material is not included in the chapter's Creative Commons license and your intended use is not permitted by statutory regulation or exceeds the permitted use, you will need to obtain permission directly from the copyright holder.

